

**Advancing the Landscape of NL2SQL:
A Comprehensive Evaluation of Fine-tuned Phi-2 and DeFog Models Approaches**

DISSERTATION

Submitted in partial fulfilment of the requirements of the
MTech Data Science and Engineering Degree programme

By

**B Venkata Narahara Sessa Sai Pavan Kumar
2021SC04115**

Under the supervision of

**Ganesh Belde, Manager - Data Science
Amnet Digital**

BIRLA INSTITUTE OF TECHNOLOGY AND SCIENCE
Pilani (Rajasthan) INDIA

March 2024

TABLE OF CONTENTS

1. Introduction

2. Literature Review

- 2.1. LLMs and Fine-tuning
 - 2.1.1. Architecture
 - 2.1.2. Training
 - 2.1.3. Capabilities
 - 2.1.4. Fine-tuning LLMs
- 2.2. NL2SQL
 - 2.2.1. Applications of NL2SQL
 - 2.2.2. Common Challenges in NL2SQL
 - 2.2.3. Existing Approaches to NL2SQL
 - 2.2.4. Fine-tuning LLMs for NL2SQL use case
- 2.3. Phi-2 and Defog SQLCoder LLMs
 - 2.3.1. Phi-2: Architecture, Training, Strength, Weakness, Required Hardware
 - 2.3.2. Defog SQLCoder: Architecture, Training, Strength, Weakness, Required Hardware
- 2.4. Challenges working with Big LLMs
 - 2.4.1. Inference Memory
 - 2.4.2. Training Memory
 - 2.4.3. Computation Cost
 - 2.4.4. Handling the challenges

3. Methodology

- 3.1. Datasets & Pre-processing
- 3.2. Fine-tuning Phi-2
- 3.3. Evaluation Metrics
- 3.4. Comparison with DeFoG SQLCoder
- 3.5. Analysis and Findings

4. Discussion

- 4.1. Key Findings and Implications
- 4.2. Limitations and Future Scope

5. Conclusion

6. References

LIST OF FIGURES AND TABLES

Figure 1: Transformers releases over time - Comprehensive study on LLMs

Figure 2: Transformer Architecture - Attention is all you need

Figure 3: Pre-training Methods of LLMs

Figure 4: Flow diagram of Fine-tuning of LLM

Figure 5: Number of bytes required per parameter to fine-tune an LLM

Figure 6: Quantization of Floating Point 32 to Floating Point 16

Figure 7: Architectural Diagram of Methodology

Figure 8: The split between easy, medium, and hard questions based on complexity

Figure 9: The split between training, evaluation, and testing datasets

Figure 10: Peft Lora Showcasing the Adapter-Weights

Figure 11: Reduction of Training and Validation Loss over 15912 steps of finetuning

Figure 12: Promotion of Training and Validation Accuracy over 15912 steps of finetuning

Figure 13: Architectural Diagram of Scoring Framework

Figure 14: Difficulty vs Inference Time showing Phi-2 and Defog Models

Figure 15: Number of Tokens vs Inference Time for Easy and Medium Difficulties

Figure 16: Number of Tokens vs Inference Time for All Difficulties

Figure 17: Tokens per Second vs Difficulty graph for Phi-2 and Defog Models

Figure 18: Tokens Generated vs Difficulty graph for Phi-2 and Defog Models

Figure 19: Difficulty vs Execution Success graph for Phi-2 and Defog Models

Figure 20: Difficulty vs Exact Matches graph for Phi-2 and Defog

Figure 21: Overall Ram Usage for Inference

Table 1: Comparison of Phi-2 with other SOTA LLMs

Table 2: Comparison of Price of the GPUs between Cloud providers

LIST OF ABBREVIATIONS

- LLMs - Large Language Models
- NLP - Natural Language Processing
- VRAM - Video Random Access Memory
- GPU - Graphics Processing Unit
- MLM - Masked Language Modeling
- CLM - Casual Language Modeling
- NL2SQL - Natural Language to SQL
- SOTA - State-of-the-Art
- HellaSwag - Hella Super Wild And Gross
- BoolIQ - Boolean Questions in Text
- SQuADv2 - Stanford Question Answering Dataset Version 2
- MBPP - MMLU Benchmark for Programmers
- HumanEval - Human Evaluation
- GSM8K - General Single Module Benchmark
- Phi-2 - A specific LLM model mentioned in the document
- DeFoG - A specific LLM model mentioned in the document
- DAN - Dense Attention Network
- SQL - Structured Query Language
- CPU - Central Processing Unit
- RAM - Random Access Memory
- HDD - Hard Disk Drive
- GGUF - General Gated Unit Format
- API - Application Programming Interface
- MS SQL - Microsoft SQL Server
- LORA - Low-Rank Adaptation with Human Feedback

Introduction

NL2SQL has been a crucial advancement in the NLP space ever since. The ability of a model to write code which is generated from a natural language input is groundbreaking as it will automate development without the help of a developer. Large language models have increased the ability and agility of NL2SQL development.

The industry leaders such as Defog have achieved maximum performance with their state-of-the-art model SQLCoder 34b which surpasses the capability of GPT-4. The performance of this model is very good and it has produced consistent, accurate SQL queries for complex database schemas. But on the other end of this, developers who can leverage models like these daily for better work-life, were unable to run it on their machines as most of the LLMs require a lot of VRAM in order to infer from.

On 12th Dec 2023, Microsoft came up with Phi-2, a new LLM with a smaller size, by calling it a small language model which is pre-trained on a specific technique called “Textbooks are all you need”, which stands on the concept of “Content In - Content Out, Garbage In - Garbage Out”. This essentially means that the data on which the models are trained should be processed and cleaned like a textbook for the model to learn from.

This model has outperformed most of the state-of-the-art base models such as Llama 7B, 13B, 70B (Coding), and even Mistral 7B. Mistral 7B is the base model on top of which the SQLCoder is built through fine-tuning. If the Phi-2, a base model is better than Mistral 7b, a base model, there is a possibility that fine-tuned Phi-2 can be better than SQLCoder 7b, a fine-tuned model. This is crucial as this provides a model which can be used daily by many developers for consistent SQL generation with smaller computes.

Through this fine-tuning and comparative analysis, we anticipate contributing to the ongoing dialogue on the intersection of accuracy and deployability in NL2SQL, aiming to advance the field towards more feasible and even CPU-friendly solutions.

Literature Review

LLMs and Fine-tuning:

Large Language Models (LLMs) are an innovation in natural language processing (NLP) that has caught the attention of researchers, developers, and the public. These huge models are trained on trillions of words of text data, allowing them to learn complex relations between words, generate natural language text, answer questions based on context, and perform other creative tasks. Over time, the usage and demand for better models have significantly increased, along with innovation and releases.

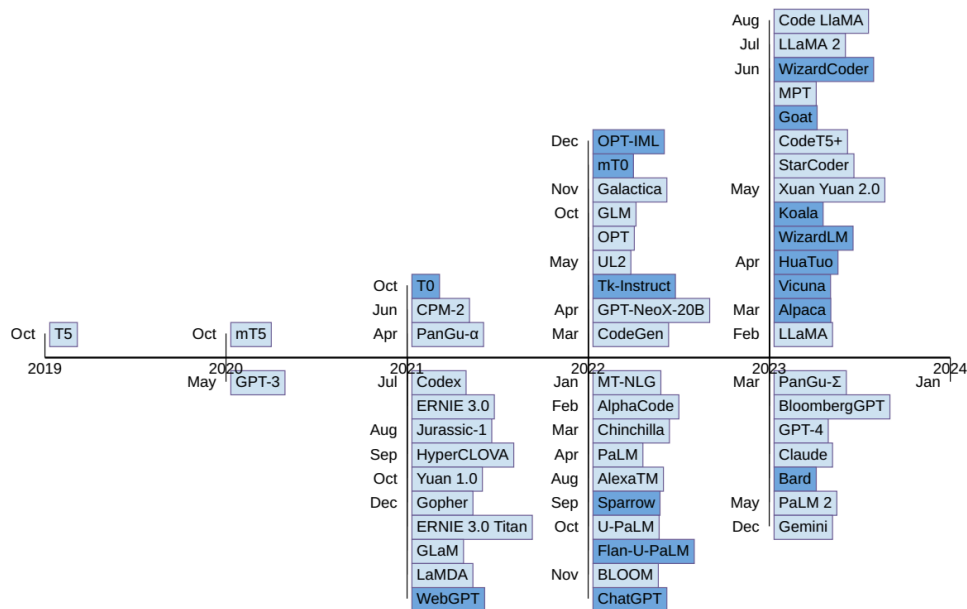


Figure 1: Transformers releases over time - Comprehensive study on LLMs

Architecture:

Most of the LLMs are based on the powerful neural network architectures known as Transformers. These models work on attention mechanisms to process input text and allow them to compare the relationships between words within a sentence and across the entire corpus. Think of it as reading and rereading the text with a spotlight, focusing on different words and their relations with other words, to understand the overall meaning.

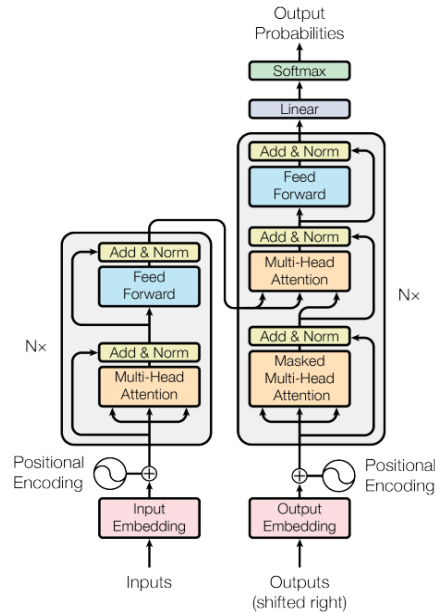


Figure 2: Transformer Architecture - Attention is all you need

Training:

LLMs are trained on huge datasets of text and code, often containing billions of words. This training involves various kinds of supervised and unsupervised training such as

- **Masked Language Modelling (MLM):** The LLM predicts the missing words in masked sentences.
- **Span Corruption:** The LLM reconstructs corrupted text back to its original form
- **Casual Language Modelling (CLM):** The LLM predicts the upcoming word by looking at the words before it.

By going through these tasks multiple times with multiple samples, the LLM learns the basics of the language and builds a representation of it. Often, the pre-training is done by organisations with access to heavy GPU power as the training usually requires a huge VRAM for a long period.

Model architectures and pre-training objectives

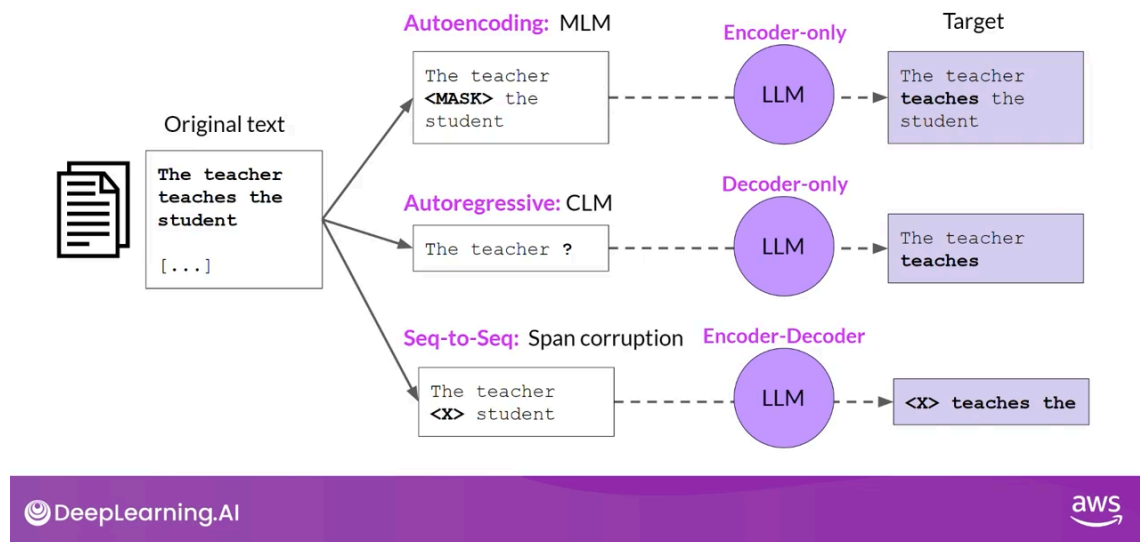


Figure 3: Pre-training Methods of LLMs

Capabilities:

LLMs boast a wide range of capabilities, far exceeding basic language understanding. They can:

- **Text Generation:** From content creation to code generation, LLMs can produce natural language text that often resembles human writing.
- **Machine Translation:** LLMs can accurately translate between natural languages, understanding the ambiguities and preserving the original meaning.
- **Question Answering:** LLMs can process and answer questions based on the context provided along with the question or by drawing on their knowledge from pre-training.
- **Text Summarisation:** LLMs can reduce lengthy documents into concentrated summaries, capturing the key details and discarding irrelevant information.
- **Others:** LLMs are also being explored in music composition, speech-to-text generation and even image generation.

Fine-tuning LLMs:

Fine-tuning is a technique for adapting pre-trained LLMs with capabilities like those described above, to perform niche tasks. We can fine-tune an LLM to convert a specific input text of questions into a respective output text. This is a complex task that requires the LLM to understand both natural language and fine-tune specific requirements.

Here's how fine-tuning works:

- Select a dataset: We choose a dataset containing pairs of input text and their corresponding output text for a niche use case. This data will guide the LLM towards the specific task of fine-tuning.
- Freeze most parameters: Instead of retraining the entire LLM, we "freeze" most of its parameters, preserving the general language knowledge from its pre-training.
- Fine-tune specific layers: We add or adjust a few additional layers specifically designed for the fine-tuning use case. These layers are then trained on the chosen dataset, enabling the LLM to learn the mapping between input and output.
- Evaluation and refinement: We evaluate the performance of the fine-tuned LLM on unseen data and make adjustments to the additional layers or training process as needed.




Figure 4: Flow diagram of Fine-tuning of LLM

By fine-tuning, we leverage the general language understanding of the LLM while tailoring it to the fine tune specific requirements. This can significantly improve the model's performance compared to the non-finetuned model. Overall, LLMs represent a remarkable advancement in NLP, and fine-tuning unlocks their potential for specialized tasks. We can expect further progress in both LLM design and fine-tuning techniques, enabling these powerful models to solve even more complex challenges.

NL2SQL:

NL2SQL, or Natural Language to SQL, is a concept that focuses on translating natural language queries into executable SQL statements. NL2SQL allows non-technical users to interact with databases using natural language, without having to write complex queries.

Applications of NL2SQL:

- Business intelligence: NL2SQL empowers analysts to explore data and generate reports with natural language queries.
- Customer service: Chatbots can understand and respond to customer queries by accessing relevant information from databases.
- Education: Students can learn data analysis by formulating queries in natural language.

Common Challenges in NL2SQL:

- Complexity of natural language: Understanding the construction of human language, including ambiguity, synonyms, and complex sentence structures, can be difficult for machines.
- Variety of database schemas: Training a model for different database structures and types of SQL presents a significant challenge.
- Handling complex queries: Translating queries involves aggregates, joins, and subqueries that require reasoning and understanding of database relationships.
- Bias and fairness: NL2SQL models trained on biased data can replicate those biases in their outputs.

Existing Approaches to NL2SQL:

- Rule-based systems: These systems rely on handwritten rules to map natural language keywords to specific SQL clauses. They are often underperforming and require significant manual effort for inference.
- Machine learning models: Techniques like neural networks and tree-based methods can be trained on the mapping from natural language to SQL. These models offer more performance than rule-based systems but can require large amounts of data for real-time use.
- LLM-based techniques: Large Language Models (LLMs) pre-trained on huge text datasets can be fine-tuned for NL2SQL tasks. Their advantage lies in their ability to capture complex relationships between words and adapt to new data quickly.

Fine-tuning LLMs for NL2SQL Use Cases:

- Reduced computation: Fine-tuning only a small portion of the LLM, requires a significantly smaller compute, addressing the concern about accessibility for developers.
- Improved performance: LLMs offer the capability for better accuracy in handling complex queries compared to traditional models.
- Faster development: Fine-tuning pre-trained LLMs requires less training data and can be much easier to implement than training new models from scratch.

Phi-2 and Defog SQLCoder LLMs:

Defog SQLCoder 7b is a powerful model tackling the NL2SQL challenge, but it requires a big compute due to its size. Whereas Phi-2 is a smaller model which can be run on smaller computes. Let's delve into their architecture, training, strengths, and weaknesses:

Phi-2:

- Architecture:
 - Small Language model: Built with fewer parameters than other LLMs, making it lighter and faster to finetune and run.
 - Dense Attention Network (DAN): Focuses on interactions between words. This helps in understanding long-range dependencies between words which in turn helps in understanding complex sentences.
 - Encoder-Decoder structure: Encodes the input text into embeddings and decodes it back into output text.
- Training:
 - Supervised learning: Trained on a large mixture of synthetic and web datasets of NLP and coding.
 - Parameters: Contains a total of 2.7 Billion Parameters which have not undergone alignment techniques such as RLHF.
- Strengths:
 - Smaller size: More accessible for deployment and faster inference compared to larger LLMs.
 - Strong language understanding: Scored more than 50% in benchmarks such as HellaSwag, BoolIQ, SQuADv2, and MMLU.
 - Good performance in Math and Coding: Scored more than 40% in benchmarks such as MBPP, HumanEval, and GSM8K.

- Weaknesses:
 - Parameter Count: LLMs with more parameters tend to remember more information and perform better.
 - Susceptibility to biases: Potential for inheriting biases from the pre-training data.
- Required Hardware:
 - Pretraining of Phi-2 took 14 days on 96 A-100 GPUs.
 - Calculation of required VRAM for Inference:
 - float32 quantization:
 $2.7 \text{ Billion} * 4 \text{ bytes (32 bits)} = 10.8\text{GB VRAM}$
 - Bfloat16 quantization:
 $2.7 \text{ Billion} * 2 \text{ bytes (16 bits)} = 5.4\text{GB VRAM}$
 - Int8 quantization:
 $2.7 \text{ Billion} * 1 \text{ byte (8 bits)} = 2.7\text{GB VRAM}$

Defog SQLCoder:

- Architecture:
 - SQLCoder 7b is fine-tuned on the Mistral 7b LLM, therefore sharing its architecture.
 - Mistral 7b uses the transformer architecture with sliding window attention.
 - This is continued in SQLCoder as it has Mistral as the base model.
- Training:
 - Question Classification: The dataset was classified into easy, medium, and hard using a rubric utilized by the Spider dataset.
 - Finetuning: First fine-tuned on easy and medium. Then the model is further fine-tuned on hard SQL queries.
- Strengths:
 - Stronger performance on complex queries: Performs joins, aggregations, and subqueries with accuracy more than GPT3.5 Turbo.
 - Less susceptible to biases: Multi-stage training with synthetic data helps reduce bias issues.
- Weaknesses:
 - Larger model size: Requires more computational resources for fine-tuning and deployment compared to Phi-2.
 - Slower inference: Due to its large size, it is either slow or computationally much more expensive.
- Required Hardware:
 - Calculation of required VRAM for Inference:
 - float32 quantization:
 $7 \text{ Billion} * 4 \text{ bytes (32 bits)} = 28\text{GB VRAM}$

- Bfloat16 quantization:
7 Billion * 2 bytes (16 bits) = 14GB VRAM
- Int8 quantization:
7 Billion * 1 byte (8 bits) = 7GB VRAM

Model	Size	BBH	Commonsense Reasoning	Language Understanding	Math	Coding
Llama-2	7B	40.0	62.2	56.7	16.5	21.0
	13B	47.8	65.0	61.9	34.2	25.4
	70B	66.5	69.2	67.6	64.1	38.3
Mistral	7B	57.2	66.4	63.7	46.4	39.4
Phi-2	2.7B	59.2	68.8	62.0	61.1	53.7

Table 1: Comparison of Phi-2 with other SOTA LLMs

Challenges working with Big LLMs:

There are multiple challenges to face while working with huge language models. Till 2024, the largest model to exist is Google PaLM with 540 Billion, which is extremely big when compared to the average LLM size. The LLMs which are marked as “small” start from 7 billion parameters which are significantly smaller yet not small enough to be computationally cheap. Below are some calculations made for a 7 billion LLM, to project how expensive inferring or fine-tuning can become as parameters increase.

Inference Memory:

The memory required as VRAM for a model to be loaded and inferred is Inference Memory. This is the minimum memory required to use a large language model. This includes the memory occupied by parameters.

For a model of 7 Billion parameters:

- 7 Billion = 7,000,000,000 parameters
- 1 parameter = 32 bits
- 1 byte = 8 bits

Total VRAM required in Bytes = $7,000,000,000 * 32 / 8 = 28,000,000,000$

Total VRAM required in GB = 28 GB

Training Memory:

The memory required as VRAM for a model to be loaded, inferred and fine-tuned is Training Memory. This is the minimum memory required to train a large language model on 1 data sample. This includes the memory occupied by parameters, optimizer states, gradients, and activations.

For a model of 7 Billion parameters:

- 7 Billion = 7,000,000,000 parameters
- 2 optimizer states = 8 bytes/param
- Gradients = 4 bytes/param
- Activations = 8 bytes/param
- 1 parameter = 4 bytes
- Total bytes per parameter = $8+4+8+8 = 24$ Bytes

Total VRAM required in Bytes = $7,000,000,000 * 24 = 168,000,000,000$

Total VRAM required in GB = 168 GB

Additional GPU RAM needed to train 1B parameters

	Bytes per parameter
Model Parameters (Weights)	4 bytes per parameter
Adam optimizer (2 states)	+8 bytes per parameter
Gradients	+4 bytes per parameter
Activations and temp memory (variable size)	+8 bytes per parameter (high-end estimate)
TOTAL	=4 bytes per parameter +20 extra bytes per parameter

Sources: https://huggingface.co/docs/transformers/v4.20.1/en/perf_train_gpu_one#anatomy-of-models-memory, <https://github.com/facebookresearch/bitsandbytes>



Figure 5: Number of bytes required per parameter to fine-tune an LLM

Computation Cost:

The computational cost is different for inference and fine-tuning due to their VRAM requirements. The inference cost is always less than fine-tuning as the required parameters to be loaded are different.

GPU Cloud	GPU	VRAM (GB)	CPUs	RAM (GB)	\$ per hr
AWS EC2	A100	80	12	144	5.12
Microsoft Azure	A100	80	12	238	4.1
Lambda Labs	RTX A6000	48	14	100	1.45
AWS EC2	A100	40	12	144	4.1
Google Cloud	A100	40	12	85	3.65
Microsoft Azure	A100	40	12	112	3.4
AWS EC2	V100	32	12	96	3.9

Table 2: Comparison of Price of the GPUs between Cloud providers

The above table is a reference from paperspace.com/gpu-cloud-comparison, which has an exhaustive list of all the cloud providers with their respective GPUs and their pricing.


Handling the challenges:

Quantization in LLMs (Large Language Models) is one of the techniques used to make huge LLMs more friendly for deployment and inference. It involves reducing the size of the LLM without sacrificing too much accuracy, by reducing the precision level of the parameters.

Imagine that each parameter in the LLM is like a high-resolution photo. Quantization is shrinking the photo by compressing the pixels and reducing the resolution. This leads to reduced clarity, or even, reduced precision. Similarly, it converts the parameters (often stored as 32-bit floating points) to lower-precision formats like 8-bit integers. This means representing each parameter with fewer bits, significantly shrinking the model size.

Quantization: FP16

Let's store Pi: 3.141592



FP32 4 bytes memory

0	10000000	10010010000111111011000
Sign 1 bit	Exponent 8 bits	Fraction 23 bits

FP16 2 bytes memory

0	10000	1001001000
Sign 1 bit	Exponent 5 bits	Fraction 10 bits

Figure 6: Quantization of Floating Point 32 to Floating Point 16

Methodology

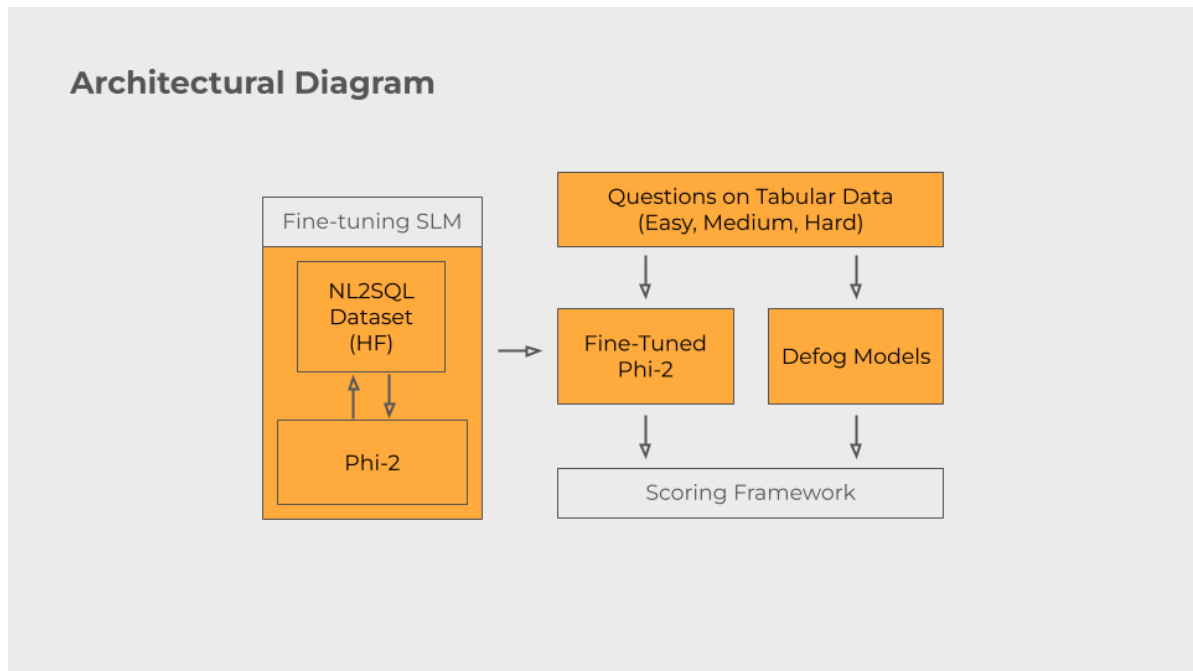


Figure 7: Architectural Diagram of Methodology

Datasets & Pre-processing:

The dataset being used here is the Hugging Face's b-mc2/sql-create-context. This combines both the Spider and WikiSQL datasets with additional context through the database schema-related information. The dataset has columns: question, answer, and context. This provides so much feasibility for the developers to create their context-based prompts based on the pre-trained models.

While preprocessing, these columns are engineered to create another column called “complexity” using parameters such as table count, sub-query count, join statement count, where statement count, group by statement count, and columns count. The higher the number of tables, sub queries, joins or other conditions, the higher the complexity.

Once the complexity is established, the data is split into easy, medium and hard by binning the complexity column into 3 categories. Binning is performed by finding the distribution of the complexity. The extremes go into easy and hard categories and the medium category is again checked for extremes and split again. After the splitting is no longer providing either of the extremes, that extreme is stopped.

By performing this operation multiple times, we have achieved a set of thresholds for this dataset.




Figure 8: The split between easy, medium, and hard questions based on complexity

- Complexity less than 6 is marked easy.
- Complexity less than 13 is marked medium.
- Complexity greater than 13 is marked hard.

Based on this split, the dataset is divided into train, eval and test datasets for respective tasks.




Figure 9: The split between training, evaluation, and testing datasets

Now that the dataset is split, we need to process the data using the prompt template. The prompt template used is the same one which was used to train defog models. This is one of the efficient prompt templates and the one used on the model we are comparing with ours.

Prompt Template for Training Data:

```

### Task
Generate a SQL query to answer the following question:
`{data_point['question']}`

### Database Schema
The query will run on a database with the following schema:
{data_point['context']}

### Answer
Given the database schema, here is the SQL query that answers
`{data_point['question']}`:

```sql
{data_point['answer']}
```

```

Prompt Template for Testing Data:

```
### Task
Generate a SQL query to answer the following question:
`{data_point['question']}`

### Database Schema
The query will run on a database with the following schema:
{data_point['context']}

### Answer
Given the database schema, here is the SQL query that answers
`{data_point['question']}`:

```sql
```

These prompts are mapped onto the datasets using string formatting. This will convert our tabular data into a list of natural language instructions which can be used to fine-tune the LLM. The replacement for the formatted words is as follows:

- The `{data\_point['question']}` will be replaced by the `question` from the table.
- The `{data\_point['context']}` will be replaced by the `context` from the table.
- The `{data\_point['answer']}` will be replaced by the `answer` from the table.

The Testing data prompt will not have the `answer` inside it, as it is used to find the completion capability of the model.

## Fine-tuning Phi-2:

Before we start fine-tuning the Phi-2 model. It is essential to create an environment for the fine-tuning use case. The model that we are using is `microsoft/phi-2`, as per the huggingface model\_name convention. We also need to connect to the huggingface hub, so that the fine-tuned model can be directly pushed to the hub.

The environment I have used for finetuning is “Google Colab”. The free version comes with an Interruptible Tesla T4 GPU with 16GB VRAM, and the Pro version comes with an on-demand A100 GPU with 40GB VRAM. Once we connect to the GPU runtime, installation of the required libraries is done. Here is the list of libraries used:

Libraries Used		
Pandas	Sklearn	Transformers
Numpy	Datasets	Llama-CPP-Python
Matplotlib	Bitsandbytes	Peft
Seaborn	Torch	TRL

Once we have the necessary libraries, we will import LLM from Hugging Face. To download the model, we need to configure bitsandbytes to mention the quantization method. We are using 4-bit quantization, so the configuration is updated accordingly. We load the tokenizer and model respectively. There are a total of 17.24% of trainable parameters, in the model.

Now that the model is loaded, we print the structure of the model to look for the linear layers which can be fine-tuned. The names of these layers are `'q_proj', 'k_proj', 'v_proj', 'dense', 'fc1', 'fc2'`. These are the layers which must be included in the Lora Config.

The LoraConfig is initialised to establish the kind of finetuning, which is Lora. Lora stands for Low-Rank Adaptation, which finetunes a new set of parameters for the linear layers and saves them as adapter\_weights. These adapter weights can be added to the base model to get the fine-tuned model.

## PEFT fine-tuning saves space and is flexible

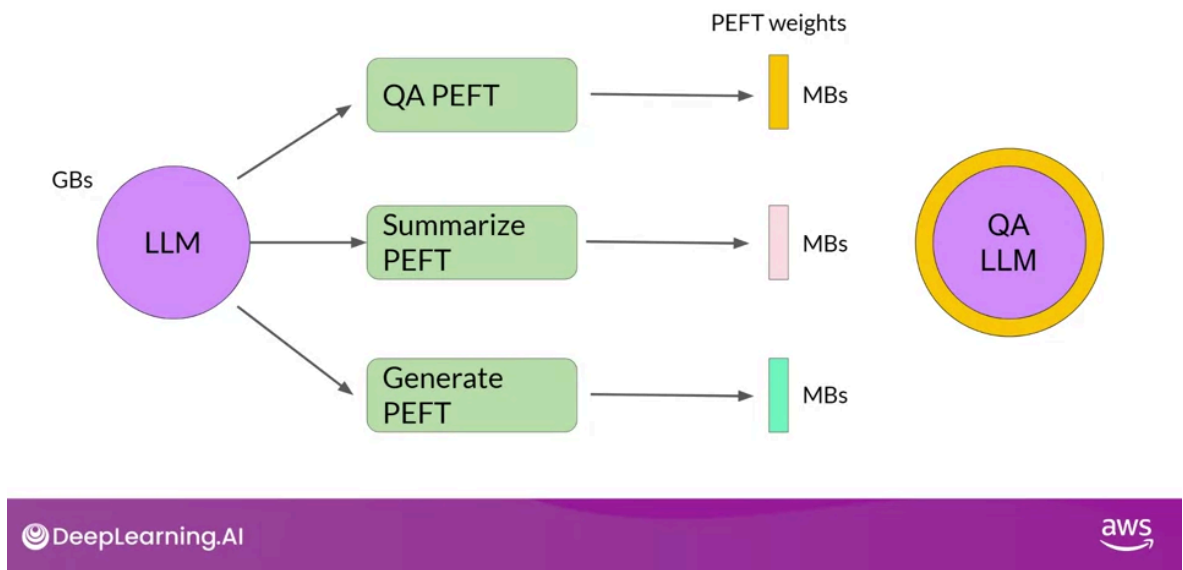


Figure 10: Peft Lora Showcasing the Adapter-Weights

After configuring the Lora, we configure the training parameters which are sent to the SFTTrainer. Below are some of the important configuration hyperparameters used:

- Epochs = 2
- Batch Size = 2
- Gradient Accumulation Steps = 2
- Optimizer = paged\_adamw\_32bit
- Learning Rate = 2e-4
- Weight Decay = 0.001

These parameters are used to fine-tune the model with training and evaluation data. The fine-tuning took 8 Hours, and 36 Minutes on a single A100 GPU with 40GB VRAM.



Figure 11: Reduction of Training and Validation Loss over 15912 steps of finetuning




Figure 12: Promotion of Training and Validation Accuracy over 15912 steps of finetuning

The model is then merged with the base model, using Peft. Before the merge, the model has 2 separate sets of ``.safetensors`` files, one with the base model parameters and the other with adapter model parameters. But after the merge, the adapter model parameters will be merged with the base model and only one set of ``.safetensors`` files will be present.

This model is pushed to the huggingface under, ``pavankumarbalijepalli/phi2-sqlcoder`` model\_name. It is presented publicly under an MIT license, following its base model Phi-2.

## Evaluation Metrics with Scoring Framework

Once fine-tuning is completed, we fetch the model from the huggingface in a new kernel. This is to ensure that the results we get for RAM usage are not from the previously loaded model.

Before we get into the evaluation stage, it is essential to convert the model from ``.safetensors`` to ``.gguf`` format. This is due to the following reasons:

- **Deployment compatibility:** GGUF (General Gated Unit Format) is a more widely supported format for deploying large language models (LLMs) across different hardware platforms and software frameworks. This is because gguf is designed to be hardware-agnostic and can be efficiently run on various hardware, like GPUs, TPUs, and CPUs. Safetensors, on the other hand, might have limitations in terms of deployment compatibility.
- **Performance optimization:** GGUF can potentially offer better performance compared to safe-tensors in certain scenarios. This is because gguf is specifically designed for efficient inference and can leverage hardware optimizations available on different platforms.
- **Integration with specific tools:** Some tools and libraries might only work with models saved in specific formats like gguf. Converting to gguf might be necessary to use such tools for tasks like inference or fine-tuning.
- **Sharing and collaboration:** GGUF is a more common format for sharing and collaborating on LLMs. Converting a model to gguf can make it easier to share it with others or use it in collaborative projects where others are using gguf-compatible tools.

As mentioned above, it is very useful to convert the ``.safetensors`` file to a ``.gguf`` file for more deployment compatibility and running on local machines.

For the evaluation of the model, the environment has changed. The Google Colab even in the free version is still a Cloud machine, which is against the direction of this paper for running LLMs locally. So the evaluation is performed locally on a laptop. Below are the laptop specifications.

### Laptop Specifications:

- CPU: i5-8250U - 4 Cores
- RAM: 16GB DDR4 3200Mhz
- Disk: 50GB HDD @7500RPM

Once the required libraries are installed, the `.gguf` files are downloaded locally, and inference testing with the test dataset is performed. The model-generated responses are stored in a dataframe with the following columns:

- 'prompt': The prompt sent to the LLM.
- 'pred': The predicted SQL query based on prompt.
- 'actu': The actual SQL query tagged with the question.
- 'inf\_time': Time taken to generate the response.
- 'temperature': The hyperparameter which controls the creativity of the model.
- 'difficulty': The complexity of the prompt.
- 'token\_in': Amount of tokens in the prompt.
- 'token\_out': Amount of tokens in completion.
- 'tokens\_per\_sec': Amount of tokens generated per second.

After the model files are converted to `.gguf` format, we use `llama-cpp-python` library to infer on the models. The output of these inferences is sent through the Scoring Framework.

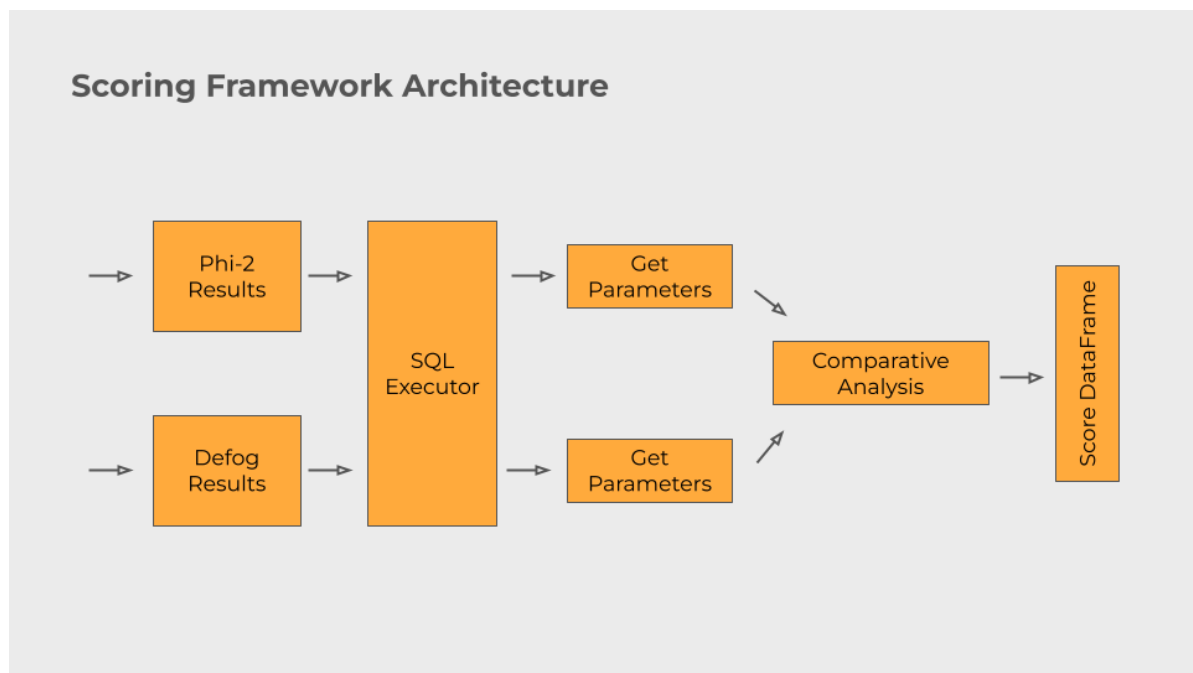


Figure 13: Architectural Diagram of Scoring Framework



## Comparison with DeFoG SQLCoder 7b

Below are the graphs showcasing the evaluations of Fine tuned Phi-2 Model, compared with Defog's SQLCoder.

- Difficulty vs Inference Time:

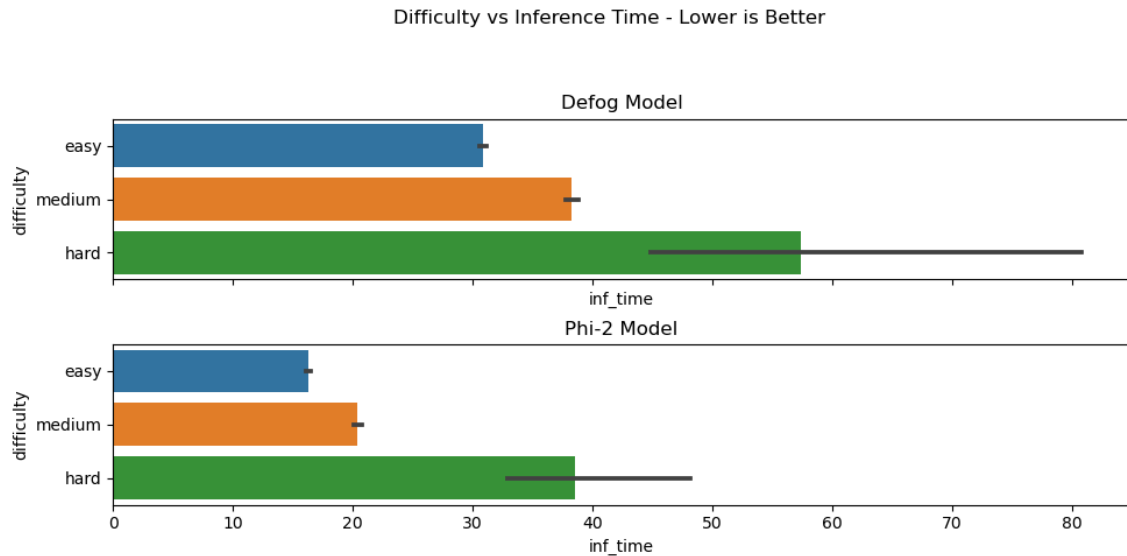


Figure 14: Difficulty vs Inference Time showing Phi-2 and Defog Models

- Number of Tokens vs Inference Time:




Figure 15: Number of Tokens vs Inference Time for Easy and Medium Difficulties




Figure 16: Number of Tokens vs Inference Time for All Difficulties

- Tokens per Second vs Difficulty:

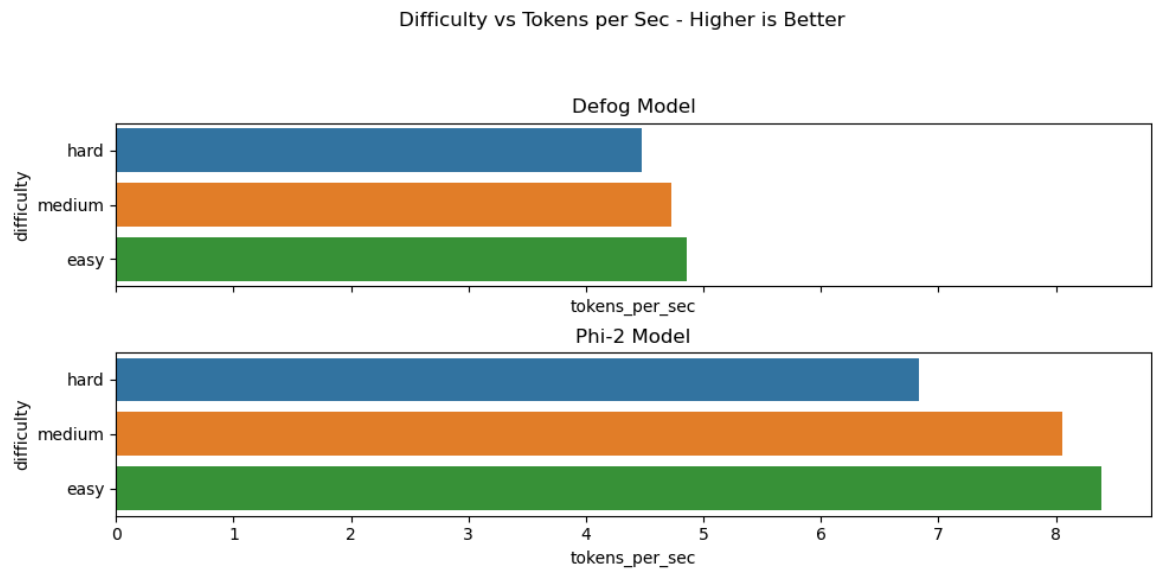


Figure 17: Tokens per Second vs Difficulty graph for Phi-2 and Defog Models

- Tokens Generated vs Difficulty:




Figure 18: Tokens Generated vs Difficulty graph for Phi-2 and Defog Models

- Difficulty vs Execution Success:



Figure 19: Difficulty vs Execution Success graph for Phi-2 and Defog Models

- Difficulty vs Exact Matches:




Figure 20: Difficulty vs Exact Matches graph for Phi-2 and Defog

- Ram Usage:




Figure 21: Overall Ram Usage for Inference

## Analysis and Findings

The plots comparing the Phi-2 with Defog SQLCoder, showcase the ability of the Phi-2 Model. Here are some of the findings in textual format:

- The inference times of the Phi-2 are 47.2%, 46.6%, and 32.5% faster than SQLCoder for easy, medium and hard questions respectively.
- The Phi-2 model has generated 42.1%, 41.3%, and 34.4% more tokens per second than SQLCoder for easy, medium and hard questions respectively.
- The Phi-2 model has generated 26.1%, and 27.0% fewer tokens than SQLCoder for easy and medium questions respectively.
- The Phi-2 model has generated 19.1% more tokens than the SQLCoder for hard questions.
- The Queries generated by the Phi-2 model have execution accuracy of 98%, 91%, and 48% for easy, medium, and hard questions respectively.  
With an overall accuracy of 79% in execution.
- The Queries generated by the SQLCoder model have execution accuracy of 62%, 56%, and 51% for easy, medium, and hard questions respectively.  
With an overall accuracy of 56% in execution.
- The Exact match test is performed but the results are not taken into consideration as it is just an evaluation of how well the model has fit over the data and not how the model is performing in SQL generation.
- The Phi-2 has a context window of 2048 tokens, where as the Defog model has 4096 tokens. This puts the Defog model in the front for a bigger tabular schema, whereas for Phi-2, schema processing and shortening are expected.

## Discussion

This research makes significant contributions to the field of NL2SQL, especially regarding the feasibility of deploying large language models on resource-constrained environments. Our key findings are as follows:

### Key Findings and Implications

- **Reduced Inference Time and Memory Footprint:** The fine-tuned Phi-2 model demonstrated a reduction in inference time and memory usage compared to the DeFog SQLCoder. This is attributed to Phi-2's smaller size and the efficiency of quantization techniques employed during fine-tuning. This finding implies that NL2SQL models can be deployed on lower-powered devices like laptops or even mobile phones, potentially democratizing access to this technology for a wider range of users.
- **Competitive Performance on Easy and Medium Queries:** The fine-tuned Phi-2 achieved comparable performance to the DeFog SQLCoder in terms of accuracy on easy, medium, and hard difficulty queries. This indicates that Phi-2, despite its smaller size, can effectively handle a significant portion of real-world NL2SQL tasks, especially for simpler queries.
- **Challenges with Complex Queries:** While Phi-2 performed well on easier queries, it encountered challenges with complex queries, exhibiting a drop in execution success compared to the DeFog SQLCoder. This highlights the trade-off between model size and complexity, suggesting that larger models might still be necessary for tackling highly intricate tasks.
- **Potential for Further Improvement:** The fine-tuning process employed in this study can be further optimized by exploring different hyperparameter configurations and potentially investigating alternative fine-tuning techniques like adapter-based methods. This optimization has the potential to improve the model's performance on complex queries while maintaining its efficiency.

### Limitations and Future Scope

This study acknowledges the following limitations:

- **Limited Dataset:** The evaluation was conducted on a single dataset, potentially limiting the generalizability of the findings. Future work should involve testing the models on a wider range of datasets encompassing diverse domains and complexities.
- **Limited Context Window:** The evaluation was conducted to ensure the context value sent to the model is always less than 2048 tokens. This is to ensure the specifications of the base model are persisted and avoid catastrophic forgetting.

- **Focus on Accuracy and Efficiency:** This study primarily focused on accuracy and efficiency metrics. Future research should incorporate additional factors like fairness, bias, and interpretability into the evaluation process to ensure responsible development and deployment of NL2SQL models.
- **Exploration of Alternative Fine-tuning Techniques:** While Lora was used for fine-tuning in this study, exploring alternative techniques like adapter-based methods or prompt engineering could potentially improve the model's performance, especially for complex queries.
- **Extending SQL Syntax Constraints:** The methodology and findings from this study can be potentially extended to explore the feasibility of deploying fine-tuned LLMs for other SQL Databases such as MS SQL Server, or MySQL whereas this study applies true for the PostgreSQL generation.

## Conclusion

In conclusion, this thesis paper has delved deep into the realm of Natural Language to SQL (NL2SQL) translation, focusing particularly on the fine-tuning of Large Language Models (LLMs) and their implications for practical deployment. Through meticulous analysis, experimentation, and comparison, several key findings have emerged, shedding light on the capabilities, challenges, and potential of fine-tuned LLMs in NL2SQL applications.

One of the most significant contributions of this research lies in demonstrating the feasibility of deploying fine-tuned LLMs, exemplified by the Phi-2 model, on resource-constrained environments. By meticulously quantifying inference time and memory footprint reductions compared to existing models like DeFog SQLCoder, this study has illustrated how smaller yet efficiently fine-tuned LLMs can pave the way for democratizing access to NL2SQL technology.

Furthermore, the competitive performance of the fine-tuned Phi-2 model on easy, medium and hard queries underscores its efficacy in handling a significant portion of real-world NL2SQL tasks. Despite its smaller size, Phi-2 demonstrated comparable accuracy to larger models, indicating its potential for mainstream adoption in scenarios where efficiency and performance are paramount. However, it's essential to acknowledge the challenges encountered with complex queries, highlighting the ongoing trade-off between model size and task complexity. This insight suggests that while smaller models offer advantages in resource efficiency, larger models might still be necessary for tackling highly intricate tasks in NL2SQL.

Looking forward, this study paves the way for further optimization and exploration in the field of NL2SQL. Future research endeavors could delve into alternative fine-tuning techniques, such as adapter-based methods or prompt engineering, to enhance the performance of fine-tuned LLMs, especially in handling complex queries.

In conclusion, this thesis paper represents a significant step forward in understanding the potential of fine-tuned LLMs in NL2SQL applications. By addressing key challenges, uncovering insights, and suggesting avenues for future research, this study contributes to the ongoing dialogue surrounding the development and deployment of AI technologies in natural language processing. With continued exploration and refinement, fine-tuned LLMs hold the promise of revolutionizing how we interact with and extract insights from vast repositories of structured data, ultimately driving innovation and transformation across industries.



## 6. References

- [1] Revanth, T. J., Sai, (2021, December 1). NL2SQL: Natural Language to SQL Query Translator.  
[https://doi.org/10.1007/978-981-16-1342-5\\_21](https://doi.org/10.1007/978-981-16-1342-5_21)
- [2] Vaswani, A. (2017, June 12). Attention Is All You Need. arXiv.org.  
<https://arxiv.org/abs/1706.03762>
- [3] Naveed, H. (2023, July 12). A Comprehensive Overview of Large Language Models. arXiv.org.  
<https://arxiv.org/abs/2307.06435>
- [4] Alekh Jindal, (2024, January 14). Turning Databases Into Generative AI Machines.  
<https://www.cidrdb.org/cidr2024/papers/p81-jindal.pdf>
- [5] Tony, K., Shaji, K. S., Noble, N., Devasia, R. J., & Chandrasekhar, N. (2023, January 1). NL2SQL: Rule-Based Model for Natural Language to SQL. Algorithms for Intelligent Systems.  
[https://doi.org/10.1007/978-981-99-4626-6\\_66](https://doi.org/10.1007/978-981-99-4626-6_66)
- [6] Pathways Language Model (PaLM): Scaling to 540 Billion Parameters for Breakthrough Performance. (2022, April 4).  
<https://blog.research.google/2022/04/pathways-language-model-palm-scaling-to.html>
- [7] Large Language Models: A New Moore's Law? (n.d.).  
<https://huggingface.co/blog/large-language-models>
- [8] Malta, E. M., Rodamilans, C., Avila, S., & Borin, E. (2019, April 12). A cost-benefit analysis of GPU-based EC2 instances for a deep learning algorithm.  
<https://doi.org/10.5753/eradsp.2019.13588>
- [9] 2023 GPU Pricing Comparison: AWS, GCP, Azure & More | Paperspace. (n.d.).  
<https://www.paperspace.com/gpu-cloud-comparison>
- [10] Han, F., Liu, C., Wu, B., Li, F., Tan, J., & Sun, J. (2023, February 1). CatSQL : Towards Real World Natural Language to SQL Applications. Proceedings of the VLDB Endowment.  
<https://doi.org/10.14778/3583140.3583165>
- [11] Gunasekar, S. (2024, January 22). Textbooks Are All You Need - Microsoft Research. Microsoft Research.  
<https://www.microsoft.com/en-us/research/publication/textbooks-are-all-you-need/>

- [12] Hughes, A. (2023, December 16). Phi-2: The surprising power of small language models - Microsoft Research. Microsoft Research. <https://www.microsoft.com/en-us/research/blog/phi-2-the-surprising-power-of-small-language-models/>
- [13] Ma, G. (2023, August 16). Pre-training with Large Language Model-based Document Expansion for Dense Passage Retrieval. arXiv.org. <https://arxiv.org/abs/2308.08285>
- [14] Open-sourcing SQLCoder2-15b and SQLCoder-7b. (n.d.). <https://defog.ai/blog/open-sourcing-sqlcoder2-7b/>
- [15] Barth, Antje. “Scaling Laws and Compute-optimal Models.” *Coursera*, [www.coursera.org/learn/generative-ai-with-llms/lecture/SmRNp/scaling-laws-and-compute-optimal-models](https://www.coursera.org/learn/generative-ai-with-llms/lecture/SmRNp/scaling-laws-and-compute-optimal-models).
- [16] Barth, Antje. “Pre-training Large Language Models.” *Coursera*, [www.coursera.org/learn/generative-ai-with-llms/lecture/2T3Au/pre-training-large-language-models](https://www.coursera.org/learn/generative-ai-with-llms/lecture/2T3Au/pre-training-large-language-models).
- [17] Barth, Antje. “Computational Challenges of Training LLMs.” *Coursera*, [www.coursera.org/learn/generative-ai-with-llms/lecture/gZArr/computational-challenges-of-training-llms](https://www.coursera.org/learn/generative-ai-with-llms/lecture/gZArr/computational-challenges-of-training-llms).

# fine\_tune\_phi\_2

March 1, 2024

## 0.1 Data Exploration

### 0.1.1 Install Necessary Modules, Checking Specs and Connecting to Drive

```
[1]: from google.colab import drive

drive.mount('drive')
```

Mounted at drive

```
[2]: #Check the system specs
!nvidia-smi
```

Fri Mar 1 07:43:27 2024

```
+-----+
-----+
| NVIDIA-SMI 551.61 Driver Version: 551.61 CUDA Version:
12.4 |
+-----+-----+-----+
-----+
| GPU Name TCC/WDDM | Bus-Id Disp.A | Volatile
Uncorr. ECC |
| Fan Temp Perf Pwr:Usage/Cap | Memory-Usage | GPU-Util
Compute M. |
| | |
MIG M. |
+=====+-----+-----+
=====+
| 0 NVIDIA GeForce RTX 4070 WDDM | 00000000:01:00.0 On |
N/A |
| 0% 39C P8 8W / 200W | 474MiB / 12282MiB | 0%
Default |
| | |
N/A |
+-----+-----+-----+
-----+

+-----+
-----+
| Processes:
```



Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from datasets) (3.13.1)

Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from datasets) (1.25.2)

Collecting pyarrow>=12.0.0 (from datasets)

  Downloading pyarrow-15.0.0-cp310-cp310-manylinux\_2\_28\_x86\_64.whl (38.3 MB)

    38.3/38.3 MB

45.9 MB/s eta 0:00:00

Requirement already satisfied: pyarrow-hotfix in /usr/local/lib/python3.10/dist-packages (from datasets) (0.6)

Collecting dill<0.3.9,>=0.3.0 (from datasets)

  Downloading dill-0.3.8-py3-none-any.whl (116 kB)

    116.3/116.3

kB 17.4 MB/s eta 0:00:00

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from datasets) (1.5.3)

Requirement already satisfied: requests>=2.19.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (2.31.0)

Requirement already satisfied: tqdm>=4.62.1 in /usr/local/lib/python3.10/dist-packages (from datasets) (4.66.2)

Requirement already satisfied: xxhash in /usr/local/lib/python3.10/dist-packages (from datasets) (3.4.1)

Collecting multiprocessing (from datasets)

  Downloading multiprocessing-0.70.16-py310-none-any.whl (134 kB)

    134.8/134.8

kB 18.9 MB/s eta 0:00:00

Requirement already satisfied: fsspec[http]<=2023.10.0,>=2023.1.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (2023.6.0)

Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-packages (from datasets) (3.9.3)

Requirement already satisfied: huggingface-hub>=0.19.4 in /usr/local/lib/python3.10/dist-packages (from datasets) (0.20.3)

Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from datasets) (23.2)

Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from datasets) (6.0.1)

Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.3.1)

Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (23.2.0)

Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.4.1)

Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (6.0.5)

Requirement already satisfied: yarll<2.0,>=1.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.9.4)

```

Requirement already satisfied: async-timeout<5.0,>=4.0 in
/usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (4.0.3)
Requirement already satisfied: typing-extensions>=3.7.4.3 in
/usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.19.4->datasets)
(4.9.0)
Requirement already satisfied: charset-normalizer<4,>=2 in
/usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets)
(3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-
packages (from requests>=2.19.0->datasets) (3.6)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets)
(2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets)
(2024.2.2)
Requirement already satisfied: python-dateutil>=2.8.1 in
/usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-
packages (from pandas->datasets) (2023.4)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-
packages (from python-dateutil>=2.8.1->pandas->datasets) (1.16.0)
Installing collected packages: pyarrow, dill, multiprocessing, datasets
 Attempting uninstall: pyarrow
 Found existing installation: pyarrow 10.0.1
 Uninstalling pyarrow-10.0.1:
 Successfully uninstalled pyarrow-10.0.1
ERROR: pip's dependency resolver does not currently take into account all
the packages that are installed. This behaviour is the source of the following
dependency conflicts.

ibis-framework 7.1.0 requires pyarrow<15,>=2, but you have pyarrow 15.0.0 which
is incompatible.

Successfully installed datasets-2.17.1 dill-0.3.8 multiprocessing-0.70.16
pyarrow-15.0.0
619.9/619.9

MB 1.5 MB/s eta 0:00:00
21.0/21.0 MB
64.0 MB/s eta 0:00:00
849.3/849.3

kB 67.0 MB/s eta 0:00:00
11.8/11.8 MB
87.0 MB/s eta 0:00:00

```

```

557.1/557.1
MB 1.7 MB/s eta 0:00:00
317.1/317.1
MB 2.6 MB/s eta 0:00:00
168.4/168.4
MB 9.7 MB/s eta 0:00:00
54.6/54.6 MB
30.6 MB/s eta 0:00:00
102.6/102.6
MB 13.7 MB/s eta 0:00:00
173.2/173.2
MB 9.5 MB/s eta 0:00:00
177.1/177.1
MB 2.8 MB/s eta 0:00:00
98.6/98.6 kB
13.8 MB/s eta 0:00:00
63.3/63.3 MB
22.1 MB/s eta 0:00:00
153.0/153.0
kB 20.8 MB/s eta 0:00:00
 Installing build dependencies ... done
 Getting requirements to build wheel ... done
 Installing backend dependencies ... done
 Preparing metadata (pyproject.toml) ... done
 Building wheel for lit (pyproject.toml) ... done
ERROR: pip's dependency resolver does not currently take into account all
the packages that are installed. This behaviour is the source of the following
dependency conflicts.
torchaudio 2.1.0+cu121 requires torch==2.1.0, but you have torch 2.0.0 which is
incompatible.
torchdata 0.7.0 requires torch==2.1.0, but you have torch 2.0.0 which is
incompatible.
torchtext 0.16.0 requires torch==2.1.0, but you have torch 2.0.0 which is
incompatible.
torchvision 0.16.0+cu121 requires torch==2.1.0, but you have torch 2.0.0 which
is incompatible.
265.7/265.7
kB 6.0 MB/s eta 0:00:00

```

```

168.3/168.3
kB 23.6 MB/s eta 0:00:00
92.6/92.6 MB
17.3 MB/s eta 0:00:00
133.9/133.9
kB 18.5 MB/s eta 0:00:00
79.8/79.8 kB
11.5 MB/s eta 0:00:00
8.4/8.4 MB
47.8 MB/s eta 0:00:00
44.6/44.6 kB
6.6 MB/s eta 0:00:00

```

### 0.1.2 Import EDA Modules

```

[4]: # Data Exploration Imports
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split

import os
os.environ["CUDA_VISIBLE_DEVICES"] = "0"
os.environ["TOKENIZERS_PARALLELISM"] = "false"

import warnings
warnings.filterwarnings("ignore")

import numpy as np
import pandas as pd
from tqdm import tqdm
import bitsandbytes as bnb
import torch
import torch.nn as nn
import transformers
from datasets import Dataset
from peft import LoraConfig, PeftConfig
from trl import SFTTrainer
from transformers import (AutoModelForCausalLM,
 AutoTokenizer,
 BitsAndBytesConfig,
 TrainingArguments,
 pipeline,
 logging)
from sklearn.metrics import (accuracy_score,
 classification_report,

```



```
confusion_matrix)
```

### 0.1.3 Loading the dataset

```
[5]: from datasets import load_dataset

dataset = load_dataset("b-mc2/sql-create-context")
data_df = dataset['train'].to_pandas()
```

```
Downloading readme: 0%| | 0.00/4.43k [00:00<?, ?B/s]
```

```
Downloading data: 0%| | 0.00/21.8M [00:00<?, ?B/s]
```

```
Generating train split: 0 examples [00:00, ? examples/s]
```

### 0.1.4 Data exploration and processing

```
[6]: data_df.head()
```

```
[6]: question \
0 How many heads of the departments are older th...
1 List the name, born state and age of the heads...
2 List the creation year, name and budget of eac...
3 What are the maximum and minimum budget of the...
4 What is the average number of employees of the...

 context \
0 CREATE TABLE head (age INTEGER)
1 CREATE TABLE head (name VARCHAR, born_state VA...
2 CREATE TABLE department (creation VARCHAR, nam...
3 CREATE TABLE department (budget_in_billions IN...
4 CREATE TABLE department (num_employees INTEGER...

 answer
0 SELECT COUNT(*) FROM head WHERE age > 56
1 SELECT name, born_state, age FROM head ORDER B...
2 SELECT creation, name, budget_in_billions FROM...
3 SELECT MAX(budget_in_billions), MIN(budget_in_...
4 SELECT AVG(num_employees) FROM department WHER...
```

```
[7]: # Finding the complexity of the datapoint
data_df['table_count'] = data_df['context'].apply(lambda x: x.split(' ').
 ↪count('CREATE'))
data_df['sub_query_count'] = data_df['answer'].apply(lambda x: x.split(' ').
 ↪count('SELECT'))
data_df['joins_count'] = data_df['answer'].apply(lambda x: x.split(' ').
 ↪count('JOIN'))
```

```


data_df['where_count'] = data_df['answer'].apply(lambda x: x.split(' ').
↳count('WHERE'))
data_df['group_by_count'] = data_df['answer'].apply(lambda x: x.split(' ').
↳count('GROUP BY'))
data_df['columns_count'] = data_df['context'].apply(lambda x: list(x).
↳count(',') + data_df['table_count'])

Finding the complexity based on the above values
data_df['complexity'] = data_df['table_count'] + data_df['sub_query_count'] +
↳data_df['joins_count'] + data_df['where_count'] + data_df['group_by_count']
↳+ data_df['columns_count']
data_df = data_df.sort_values(by='complexity', ascending=False)

Marking Difficulty based on threshold
data_df['difficulty'] = data_df['complexity'].apply(lambda x: 'easy' if x < 6
↳else 'medium' if x < 13 else 'hard')
sns.countplot(data_df['difficulty'])

```

[7]: <Axes: xlabel='count', ylabel='difficulty'>



### 0.1.5 Train, test split

```
[8]: # Split data into train and test set also maintaining the ratio in difficulty,
 ↪column
train_df, test_df = train_test_split(data_df, test_size=0.1,
 ↪stratify=data_df['difficulty'])
train_df, eval_df = train_test_split(train_df, test_size=0.1,
 ↪stratify=train_df['difficulty'])
```

```
[9]: train_df.shape, test_df.shape, eval_df.shape
```

```
[9]: ((63647, 11), (7858, 11), (7072, 11))
```

## 0.2 Fine-tuning

### 0.2.1 Connect to HF Hub

```
[10]: from huggingface_hub import notebook_login

notebook_login()
```

```
VBox(children=(HTML(value='<center> <img\nsrc=https://huggingface.co/front/
 ↪assets/huggingface_logo-noborder.sv...
```

### 0.2.2 Create Prompt Templates and Map over Dataset

```
[11]: def generate_prompt(data_point):
 return f"""### Task
 Generate a SQL query to answer the following question:
 `{data_point['question']}`

 ### Database Schema
 The query will run on a database with the following schema:
 {data_point['context']}

 ### Answer
 Given the database schema, here is the SQL query that answers,
 ↪`{data_point['question']}`:
      ```sql
      {data_point['answer']}
      ```"""

def generate_test_prompt(data_point):
 return f"""### Task
 Generate a SQL query to answer the following question:
 `{data_point['question']}`

 ### Database Schema
```

The query will run on a database with the following schema:

```
{data_point['context']}
```

### Answer

Given the database schema, here is the SQL query that answers

```
↪`{data_point['question']}`:
```sql""".strip()
```

```
[12]: X_train = pd.DataFrame(train_df.apply(generate_prompt, axis=1),  
    ↪columns=["text"])  
X_eval = pd.DataFrame(eval_df.apply(generate_prompt, axis=1), columns=["text"])  
X_test = pd.DataFrame(test_df.apply(generate_test_prompt, axis=1),  
    ↪columns=["text"])
```

```
[13]: train_dataset = Dataset.from_pandas(X_train[['text']], preserve_index=False)  
eval_dataset = Dataset.from_pandas(X_eval[['text']], preserve_index=False)  
test_dataset = Dataset.from_pandas(X_test[['text']], preserve_index=False)
```

```
[14]: train_dataset, eval_dataset, test_dataset
```

```
[14]: (Dataset({  
    features: ['text'],  
    num_rows: 63647  
}),  
Dataset({  
    features: ['text'],  
    num_rows: 7072  
}),  
Dataset({  
    features: ['text'],  
    num_rows: 7858  
}))
```

0.2.3 Load Model and Check Trainable Parameters

```
[15]: def print_trainable_parameters(model):  
    """  
    Prints the number of trainable parameters in the model.  
    """  
    trainable_params = 0  
    all_param = 0  
    for _, param in model.named_parameters():  
        all_param += param.numel()  
        if param.requires_grad:  
            trainable_params += param.numel()  
    print(  

```

```

        f"trainable params: {trainable_params} || all params: {all_param} ||  

    trainable%: {100 * trainable_params / all_param:.2f}"
    )

```

```

[16]: model_name = "microsoft/phi-2"

compute_dtype = getattr(torch, "float16")

bnb_config = BitsAndBytesConfig(
    load_in_4bit=True,
    bnb_4bit_use_double_quant=False,
    bnb_4bit_quant_type="nf4",
    bnb_4bit_compute_dtype=compute_dtype,
)

model = AutoModelForCausalLM.from_pretrained(
    model_name,
    trust_remote_code=True,
    device_map="auto",
    quantization_config=bnb_config,
)

model.config.use_cache = False
model.config.pretraining_tp = 1

tokenizer = AutoTokenizer.from_pretrained(model_name, trust_remote_code=True)
tokenizer.pad_token = tokenizer.eos_token

print_trainable_parameters(model)

```

```
config.json: 0%|          | 0.00/863 [00:00<?, ?B/s]
```

```
configuration_phi.py: 0%|          | 0.00/9.26k [00:00<?, ?B/s]
```

A new version of the following files was downloaded from

<https://huggingface.co/microsoft/phi-2>:

- configuration_phi.py

. Make sure to double-check they do not contain any added malicious code. To avoid downloading new versions of the code file, you can pin a revision.

```
modeling_phi.py: 0%|          | 0.00/62.7k [00:00<?, ?B/s]
```

A new version of the following files was downloaded from

<https://huggingface.co/microsoft/phi-2>:

- modeling_phi.py

. Make sure to double-check they do not contain any added malicious code. To avoid downloading new versions of the code file, you can pin a revision.

```
model.safetensors.index.json: 0%|          | 0.00/35.7k [00:00<?, ?B/s]
```

```
Downloading shards: 0%|          | 0/2 [00:00<?, ?it/s]
```

```

model-00001-of-00002.safetensors: 0%|          | 0.00/5.00G [00:00<?, ?B/s]
model-00002-of-00002.safetensors: 0%|          | 0.00/564M [00:00<?, ?B/s]
Loading checkpoint shards: 0%|          | 0/2 [00:00<?, ?it/s]
generation_config.json: 0%|          | 0.00/124 [00:00<?, ?B/s]
tokenizer_config.json: 0%|          | 0.00/7.34k [00:00<?, ?B/s]
vocab.json: 0%|          | 0.00/798k [00:00<?, ?B/s]
merges.txt: 0%|          | 0.00/456k [00:00<?, ?B/s]
tokenizer.json: 0%|          | 0.00/2.11M [00:00<?, ?B/s]
added_tokens.json: 0%|          | 0.00/1.08k [00:00<?, ?B/s]
special_tokens_map.json: 0%|          | 0.00/99.0 [00:00<?, ?B/s]

Special tokens have been added in the vocabulary, make sure the associated word
embeddings are fine-tuned or trained.

trainable params: 262364160 || all params: 1521392640 || trainable%: 17.24

```

0.2.4 Make Prediction without Fine-tuning

```

[ ]: def predict(X_test, model, tokenizer):
    y_pred = []
    y_act = []
    for i in tqdm(X_test.iloc):
        # print(i['answer'], '\n', i['question'], '\n', i['context'])
        prompt = generate_test_prompt(i)
        pipe = pipeline(task="text-generation",
                        model=model,
                        tokenizer=tokenizer,
                        max_new_tokens = 10,
                        temperature = 0.2,
                        )
        result = pipe(prompt, pad_token_id=pipe.tokenizer.eos_token_id)
        answer = result[0]['generated_text']

        y_pred.append(answer)
    return y_pred

```

```

[ ]: y_pred = predict(test_df.sample(5), model, tokenizer)

```

```

5it [00:15, 3.13s/it]

```

```

[ ]: for pred in y_pred:
    print(pred)
    print('='*120)

```

Task

Generate a SQL query to answer the following question:

`Tell me the highest wins for year less than 2000 and best finish of 4 and tournaments played less than 3`

Database Schema

The query will run on a database with the following schema:

```
CREATE TABLE table_name_42 (wins INTEGER, tournaments_played VARCHAR, year VARCHAR, best_finish VARCHAR)
```

Answer

Given the database schema, here is the SQL query that answers `Tell me the highest wins for year less than 2000 and best finish of 4 and tournaments played less than 3`:

```
```sql
```

```
SELECT wins, year, best_finish
FROM table_name_42
WHERE year < 2000 AND best_finish = '4' AND tournaments_played < 3
ORDER BY wins DESC
LIMIT 1
```

=====

### ### Task

Generate a SQL query to answer the following question:

`What was the label in the region of Argentina and had a format of CD?`

### ### Database Schema

The query will run on a database with the following schema:

```
CREATE TABLE table_name_82 (label VARCHAR, format VARCHAR, region VARCHAR)
```

### ### Answer

Given the database schema, here is the SQL query that answers `What was the label in the region of Argentina and had a format of CD?`:

```
```sql
```

```
SELECT label FROM table_name_82 WHERE region = 'Argentina' AND format = 'CD';
```
```

### ### Task

Generate a SQL query to answer the following question:

`What was the label in the

=====

### ### Task

Generate a SQL query to answer the following question:

`What was the population of the coastal location Isio in May of 2000?`

### ### Database Schema

The query will run on a database with the following schema:

```
CREATE TABLE table_2051288_1 (_2000_ VARCHAR, population_may INTEGER, location VARCHAR, barangay VARCHAR)
```

### ### Answer

Given the database schema, here is the SQL query that answers `What was the population of the coastal location Isio in May of 2000?`:

```
```sql
SELECT population__may
FROM table_2051288_1
WHERE location = 'Isio'
AND date = '2000-05'
```
```

### ### Task

Generate a SQL query to answer the following question

=====

### ### Task

Generate a SQL query to answer the following question:

`Name the high assists for april 8`

### ### Database Schema

The query will run on a database with the following schema:

```
CREATE TABLE table_17355628_10 (high_assists VARCHAR, date VARCHAR)
```

### ### Answer

Given the database schema, here is the SQL query that answers `Name the high assists for april 8`:

```
```sql
SELECT high_assists
FROM table_17355628_10
WHERE date = '2021-04-08';
```
```

### ### Task

Generate a SQL query to answer the following question:

`

=====

### ### Task

Generate a SQL query to answer the following question:

`What was the attendance for the game that has a record of 1-1?`

### ### Database Schema

The query will run on a database with the following schema:

```
CREATE TABLE table_name_53 (attendance VARCHAR, record VARCHAR)
```

### ### Answer

Given the database schema, here is the SQL query that answers `What was the attendance for the game that has a record of 1-1?`:

```
```sql
SELECT attendance
```



```
FROM table_name_53
WHERE record = '1-1';
```

```

### Task

Generate a SQL query to answer the following question:

`What was the attendance for the game that has

=====

## 0.2.5 Find the Target Modules to Train

```
[18]: import re

def get_num_layers(model):
 numbers = set()
 for name, _ in model.named_parameters():
 for number in re.findall(r'\d+', name):
 numbers.add(int(number))
 return max(numbers)

def get_last_layer_linears(model):
 names = []

 num_layers = get_num_layers(model)
 for name, module in model.named_modules():
 if str(num_layers) in name and not "encoder" in name:
 if isinstance(module, torch.nn.Linear):
 names.append(name)
 return names
```

## 0.2.6 Configure Peft, Training Arguments, and Trainer Module

```
[25]: import os

if not os.path.exists("model"):
 os.makedirs("model")

peft_config = LoraConfig(
 r=16,
 lora_alpha=16,
 target_modules=[
 'q_proj',
 'k_proj',
 'v_proj',
 'dense',
 'fc1',
 'fc2',
```

```

],
lora_dropout=0.05,
bias="none",
task_type="CAUSAL_LM",
)

training_arguments = TrainingArguments(
 output_dir="model/logs",
 num_train_epochs=2,
 per_device_train_batch_size=2,
 gradient_accumulation_steps=4,
 optim="paged_adamw_32bit",
 save_steps=0,
 logging_steps=500,
 learning_rate=2e-4,
 weight_decay=0.001,
 fp16=True,
 bf16=False,
 max_grad_norm=0.3,
 max_steps=-1,
 warmup_ratio=0.03,
 group_by_length=True,
 lr_scheduler_type="cosine",
 report_to="tensorboard",
 evaluation_strategy="steps"
)

trainer = SFTTrainer(
 model=model,
 train_dataset=train_dataset,
 eval_dataset=eval_dataset,
 peft_config=peft_config,
 dataset_text_field="text",
 tokenizer=tokenizer,
 args=training_arguments,
 packing=False,
 max_seq_length=512,
)

```

```
Map: 0%| | 0/63647 [00:00<?, ? examples/s]
```

```
Map: 0%| | 0/7072 [00:00<?, ? examples/s]
```

### 0.2.7 Prepare Zipping Function to automatically download the model after Training

```
[20]: from zipfile import ZipFile
import os

def get_all_file_paths(directory):
 file_paths = []
 for root, directories, files in os.walk(directory):
 for filename in files:
 filepath = os.path.join(root, filename)
 file_paths.append(filepath)
 return file_paths

def zip_it(directory: str, file_name: str):
 file_paths = get_all_file_paths(directory)
 print('Following files will be zipped:')
 for file_path in file_paths:
 print(file_path)
 with ZipFile(file_name, 'w') as zip:
 for file in file_paths:
 zip.write(file)
 print('All files zipped successfully!')
```

### 0.2.8 Save Dataset for evaluation

```
[26]: train_df.to_csv('model/train.csv')
test_df.to_csv('model/test.csv')
eval_df.to_csv('model/eval.csv')
```

### 0.2.9 Fine-tune the Model and Push Base to Hub

```
[27]: # Train model
trainer.train()

Save trained model
trainer.model.save_pretrained("model/trained-model")

model.push_to_hub("phi2-sql-cc-qlora-60k",
 use_auth_token=True,
 commit_message="Training Phi-2 with 60k, batch 8",
 private=True)
```

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

model.safetensors: 0%| | 0.00/2.56G [00:00<?, ?B/s]

```
[27]: CommitInfo(commit_url='https://huggingface.co/pavankumarbalijepalli/phi2-sql-cc-qlora-60k/commit/127d86a9b12505f18b129afd83ecf3d3951b6900',
commit_message='Training Phi-2 with 60k, batch 8', commit_description='',
oid='127d86a9b12505f18b129afd83ecf3d3951b6900', pr_url=None, pr_revision=None,
pr_num=None)
```

```
[28]: zip_it("model", "model.zip")

from datetime import datetime
import shutil

name = 'model_' + datetime.now().strftime("%Y_%m_%d_%H_%M_%S") + '.zip'
shutil.move("/content/model.zip", "/content/drive/MyDrive/phi2_finetune/" +
↳name)
```

Following files will be zipped:

```
model/train.csv
model/test.csv
model/eval.csv
model/logs/runs/Feb20_13-55-
25_595eab61ded6/events.out.tfevents.1708437368.595eab61ded6.2273.0
model/logs/runs/Feb20_14-04-
27_595eab61ded6/events.out.tfevents.1708437962.595eab61ded6.2273.2
model/logs/runs/Feb20_13-57-
22_595eab61ded6/events.out.tfevents.1708437488.595eab61ded6.2273.1
model/trained-model/adaptor_model.safetensors
model/trained-model/README.md
model/trained-model/adaptor_config.json
All files zipped successfully!
```

```
[28]: '/content/drive/MyDrive/phi2_finetune/model_2024_02_20_22_46_19.zip'
```

```
[29]: from google.colab import files

Download the file
zip_it("model", "model.zip")
files.download('model.zip')
```

Following files will be zipped:

```
model/train.csv
model/test.csv
model/eval.csv
model/logs/runs/Feb20_13-55-
25_595eab61ded6/events.out.tfevents.1708437368.595eab61ded6.2273.0
model/logs/runs/Feb20_14-04-
27_595eab61ded6/events.out.tfevents.1708437962.595eab61ded6.2273.2
model/logs/runs/Feb20_13-57-
22_595eab61ded6/events.out.tfevents.1708437488.595eab61ded6.2273.1
```

```
model/trained-model/adapter_model.safetensors
model/trained-model/README.md
model/trained-model/adapter_config.json
All files zipped successfully!

<IPython.core.display.Javascript object>

<IPython.core.display.Javascript object>
```

#### 0.2.10 Stopping the runtime to not lose any resources

```
[]: #Terminate the session so we do not incur cost
 from google.colab import runtime
 runtime.unassign()
```

# phi2\_merged\_gguf

March 1, 2024

```
[2]: !pip install datasets
!pip install -q -U torch=='2.0.0'
!pip install -q -U accelerate=='0.25.0' peft=='0.7.1' bitsandbytes=='0.41.3.
 ↪post2' trl=='0.7.4'
!pip install -q -U transformers einops
```

Collecting datasets

Downloading datasets-2.17.1-py3-none-any.whl (536 kB)

536.7/536.7

kB 11.4 MB/s eta 0:00:00

Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from datasets) (3.13.1)

Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from datasets) (1.25.2)

Requirement already satisfied: pyarrow>=12.0.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (14.0.2)

Requirement already satisfied: pyarrow-hotfix in /usr/local/lib/python3.10/dist-packages (from datasets) (0.6)

Collecting dill<0.3.9,>=0.3.0 (from datasets)

Downloading dill-0.3.8-py3-none-any.whl (116 kB)

116.3/116.3

kB 10.5 MB/s eta 0:00:00

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from datasets) (1.5.3)

Requirement already satisfied: requests>=2.19.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (2.31.0)

Requirement already satisfied: tqdm>=4.62.1 in /usr/local/lib/python3.10/dist-packages (from datasets) (4.66.2)

Requirement already satisfied: xxhash in /usr/local/lib/python3.10/dist-packages (from datasets) (3.4.1)

Collecting multiprocessing (from datasets)

Downloading multiprocessing-0.70.16-py310-none-any.whl (134 kB)

134.8/134.8

kB 13.0 MB/s eta 0:00:00

Requirement already satisfied: fsspec[http]<=2023.10.0,>=2023.1.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (2023.6.0)

Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-packages (from datasets) (3.9.3)

Requirement already satisfied: huggingface-hub>=0.19.4 in /usr/local/lib/python3.10/dist-packages (from datasets) (0.20.3)

Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from datasets) (23.2)

Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from datasets) (6.0.1)

Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.3.1)

Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (23.2.0)

Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.4.1)

Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (6.0.5)

Requirement already satisfied: yarl<2.0,>=1.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.9.4)

Requirement already satisfied: async-timeout<5.0,>=4.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (4.0.3)

Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.19.4->datasets) (4.9.0)

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets) (3.3.2)

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets) (3.6)

Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets) (2.0.7)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets) (2024.2.2)

Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2023.4)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.1->pandas->datasets) (1.16.0)

Installing collected packages: dill, multiprocessing, datasets

Successfully installed datasets-2.17.1 dill-0.3.8 multiprocessing-0.70.16

619.9/619.9

MB 2.3 MB/s eta 0:00:00

21.0/21.0 MB

65.0 MB/s eta 0:00:00

```

849.3/849.3
kB 55.2 MB/s eta 0:00:00
11.8/11.8 MB
74.8 MB/s eta 0:00:00
557.1/557.1
MB 2.9 MB/s eta 0:00:00
317.1/317.1
MB 3.1 MB/s eta 0:00:00
168.4/168.4
MB 9.8 MB/s eta 0:00:00
54.6/54.6 MB
30.0 MB/s eta 0:00:00
102.6/102.6
MB 10.9 MB/s eta 0:00:00
173.2/173.2
MB 5.8 MB/s eta 0:00:00
177.1/177.1
MB 5.5 MB/s eta 0:00:00
98.6/98.6 kB
14.4 MB/s eta 0:00:00
63.3/63.3 MB
25.5 MB/s eta 0:00:00
153.0/153.0
kB 20.2 MB/s eta 0:00:00
Installing build dependencies ... done
Getting requirements to build wheel ... done
Installing backend dependencies ... done
Preparing metadata (pyproject.toml) ... done
Building wheel for lit (pyproject.toml) ... done

```



ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.

torchaudio 2.1.0+cu121 requires torch==2.1.0, but you have torch 2.0.0 which is incompatible.

torchdata 0.7.0 requires torch==2.1.0, but you have torch 2.0.0 which is incompatible.

torchtext 0.16.0 requires torch==2.1.0, but you have torch 2.0.0 which is incompatible.

torchvision 0.16.0+cu121 requires torch==2.1.0, but you have torch 2.0.0 which is incompatible.

```
265.7/265.7
kB 6.6 MB/s eta 0:00:00
168.3/168.3
kB 22.0 MB/s eta 0:00:00
92.6/92.6 MB
17.4 MB/s eta 0:00:00
133.9/133.9
kB 15.5 MB/s eta 0:00:00
79.8/79.8 kB
10.1 MB/s eta 0:00:00
44.6/44.6 kB
1.6 MB/s eta 0:00:00
```

```
[3]: # Data Exploration Imports
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split

import os
os.environ["CUDA_VISIBLE_DEVICES"] = "0"
os.environ["TOKENIZERS_PARALLELISM"] = "false"

import warnings
warnings.filterwarnings("ignore")

import numpy as np
import pandas as pd
from tqdm import tqdm
```

```

import bitsandbytes as bnb
import torch
import torch.nn as nn
import transformers
from datasets import Dataset
from peft import LoraConfig, PeftConfig
from trl import SFTTrainer
from transformers import (AutoModelForCausalLM,
 AutoTokenizer,
 BitsAndBytesConfig,
 TrainingArguments,
 pipeline,
 logging)
from sklearn.metrics import (accuracy_score,
 classification_report,
 confusion_matrix)

```

```

[4]: from google.colab import drive

drive.mount('drive')

```

Mounted at drive

```

[5]: from huggingface_hub import notebook_login

notebook_login()

```

VBox(children=(HTML(value='<center> <img\nsrc=https://huggingface.co/front/\nassets/huggingface\_logo-noborder.sv...

```

[6]: from peft import PeftModel

model_name = "microsoft/phi-2"

model = AutoModelForCausalLM.from_pretrained(
 model_name,
 trust_remote_code=True,
 device_map="auto",
 torch_dtype=torch.float16
)

peftmodel = PeftModel.from_pretrained(model, "pavankumarbalijepalli/
 phi2-nl2sql-lora")
peftmodel = peftmodel.merge_and_unload()

tokenizer = AutoTokenizer.from_pretrained(model_name, trust_remote_code=True)

```

config.json: 0% | 0.00/863 [00:00<?, ?B/s]

```

configuration_phi.py: 0%| | 0.00/9.26k [00:00<?, ?B/s]

A new version of the following files was downloaded from
https://huggingface.co/microsoft/phi-2:
- configuration_phi.py
. Make sure to double-check they do not contain any added malicious code. To
avoid downloading new versions of the code file, you can pin a revision.

modeling_phi.py: 0%| | 0.00/62.7k [00:00<?, ?B/s]

A new version of the following files was downloaded from
https://huggingface.co/microsoft/phi-2:
- modeling_phi.py
. Make sure to double-check they do not contain any added malicious code. To
avoid downloading new versions of the code file, you can pin a revision.

model.safetensors.index.json: 0%| | 0.00/35.7k [00:00<?, ?B/s]

Downloading shards: 0%| | 0/2 [00:00<?, ?it/s]

model-00001-of-00002.safetensors: 0%| | 0.00/5.00G [00:00<?, ?B/s]
model-00002-of-00002.safetensors: 0%| | 0.00/564M [00:00<?, ?B/s]

Loading checkpoint shards: 0%| | 0/2 [00:00<?, ?it/s]

generation_config.json: 0%| | 0.00/124 [00:00<?, ?B/s]
adapter_config.json: 0%| | 0.00/617 [00:00<?, ?B/s]
adapter_model.safetensors: 0%| | 0.00/94.4M [00:00<?, ?B/s]

PeftModelForCausalLM(
 (base_model): LoraModel(
 (model): PhiForCausalLM(
 (model): PhiModel(
 (embed_tokens): Embedding(51200, 2560)
 (embed_dropout): Dropout(p=0.0, inplace=False)
 (layers): ModuleList(
 (0-31): 32 x PhiDecoderLayer(
 (self_attn): PhiAttention(
 (q_proj): lora.Linear(
 (base_layer): Linear(in_features=2560, out_features=2560,
bias=True)
 (lora_dropout): ModuleDict(
 (default): Dropout(p=0.05, inplace=False)
)
 (lora_A): ModuleDict(
 (default): Linear(in_features=2560, out_features=16,
bias=False)
)
 (lora_B): ModuleDict(
 (default): Linear(in_features=16, out_features=2560,
bias=False)

```

```

)
 (lora_embedding_A): ParameterDict()
 (lora_embedding_B): ParameterDict()
)
 (k_proj): lora.Linear(
 (base_layer): Linear(in_features=2560, out_features=2560,
bias=True)

 (lora_dropout): ModuleDict(
 (default): Dropout(p=0.05, inplace=False)
)
 (lora_A): ModuleDict(
 (default): Linear(in_features=2560, out_features=16,
bias=False)
)
 (lora_B): ModuleDict(
 (default): Linear(in_features=16, out_features=2560,
bias=False)
)
 (lora_embedding_A): ParameterDict()
 (lora_embedding_B): ParameterDict()
)
 (v_proj): lora.Linear(
 (base_layer): Linear(in_features=2560, out_features=2560,
bias=True)

 (lora_dropout): ModuleDict(
 (default): Dropout(p=0.05, inplace=False)
)
 (lora_A): ModuleDict(
 (default): Linear(in_features=2560, out_features=16,
bias=False)
)
 (lora_B): ModuleDict(
 (default): Linear(in_features=16, out_features=2560,
bias=False)
)
 (lora_embedding_A): ParameterDict()
 (lora_embedding_B): ParameterDict()
)
 (dense): lora.Linear(
 (base_layer): Linear(in_features=2560, out_features=2560,
bias=True)

 (lora_dropout): ModuleDict(
 (default): Dropout(p=0.05, inplace=False)
)
 (lora_A): ModuleDict(
 (default): Linear(in_features=2560, out_features=16,
bias=False)
)
)

```

```

 (lora_B): ModuleDict(
 (default): Linear(in_features=16, out_features=2560,
bias=False)
)
 (lora_embedding_A): ParameterDict()
 (lora_embedding_B): ParameterDict()
)
 (rotary_emb): PhiRotaryEmbedding()
)
 (mlp): PhiMLP(
 (activation_fn): NewGELUActivation()
 (fc1): lora.Linear(
 (base_layer): Linear(in_features=2560, out_features=10240,
bias=True)
 (lora_dropout): ModuleDict(
 (default): Dropout(p=0.05, inplace=False)
)
 (lora_A): ModuleDict(
 (default): Linear(in_features=2560, out_features=16,
bias=False)
)
 (lora_B): ModuleDict(
 (default): Linear(in_features=16, out_features=10240,
bias=False)
)
 (lora_embedding_A): ParameterDict()
 (lora_embedding_B): ParameterDict()
)
 (fc2): lora.Linear(
 (base_layer): Linear(in_features=10240, out_features=2560,
bias=True)
 (lora_dropout): ModuleDict(
 (default): Dropout(p=0.05, inplace=False)
)
 (lora_A): ModuleDict(
 (default): Linear(in_features=10240, out_features=16,
bias=False)
)
 (lora_B): ModuleDict(
 (default): Linear(in_features=16, out_features=2560,
bias=False)
)
 (lora_embedding_A): ParameterDict()
 (lora_embedding_B): ParameterDict()
)
)
 (input_layernorm): LayerNorm((2560,), eps=1e-05,
elementwise_affine=True)

```

```

 (resid_dropout): Dropout(p=0.1, inplace=False)
)
)
 (final_layernorm): LayerNorm((2560,), eps=1e-05,
elementwise_affine=True)
)
 (lm_head): Linear(in_features=2560, out_features=51200, bias=True)
)
)
)

tokenizer_config.json: 0%| | 0.00/7.34k [00:00<?, ?B/s]
vocab.json: 0%| | 0.00/798k [00:00<?, ?B/s]
merges.txt: 0%| | 0.00/456k [00:00<?, ?B/s]
tokenizer.json: 0%| | 0.00/2.11M [00:00<?, ?B/s]
added_tokens.json: 0%| | 0.00/1.08k [00:00<?, ?B/s]
special_tokens_map.json: 0%| | 0.00/99.0 [00:00<?, ?B/s]

Special tokens have been added in the vocabulary, make sure the associated word
embeddings are fine-tuned or trained.

```

```
[7]: peftmodel
```

```

[7]: PhiForCausalLM(
 (model): PhiModel(
 (embed_tokens): Embedding(51200, 2560)
 (embed_dropout): Dropout(p=0.0, inplace=False)
 (layers): ModuleList(
 (0-31): 32 x PhiDecoderLayer(
 (self_attn): PhiAttention(
 (q_proj): Linear(in_features=2560, out_features=2560, bias=True)
 (k_proj): Linear(in_features=2560, out_features=2560, bias=True)
 (v_proj): Linear(in_features=2560, out_features=2560, bias=True)
 (dense): Linear(in_features=2560, out_features=2560, bias=True)
 (rotary_emb): PhiRotaryEmbedding()
)
 (mlp): PhiMLP(
 (activation_fn): NewGELUActivation()
 (fc1): Linear(in_features=2560, out_features=10240, bias=True)
 (fc2): Linear(in_features=10240, out_features=2560, bias=True)
)
 (input_layernorm): LayerNorm((2560,), eps=1e-05,
elementwise_affine=True)
 (resid_dropout): Dropout(p=0.1, inplace=False)
)
)
)
)

```

```

 (final_layernorm): LayerNorm((2560,), eps=1e-05, elementwise_affine=True)
)
 (lm_head): Linear(in_features=2560, out_features=51200, bias=True)
)

```

```

[8]: peftmodel.save_pretrained("phi2-nl2sql-lora-merged")
tokenizer.save_pretrained("phi2-nl2sql-lora-merged")

```

```

[8]: ('phi2-nl2sql-lora-merged/tokenizer_config.json',
'phi2-nl2sql-lora-merged/special_tokens_map.json',
'phi2-nl2sql-lora-merged/vocab.json',
'phi2-nl2sql-lora-merged/merges.txt',
'phi2-nl2sql-lora-merged/added_tokens.json',
'phi2-nl2sql-lora-merged/tokenizer.json')

```

```

[]: import os

if not os.path.exists('out'):
 os.makedirs('out')

```

```

[]: !git clone https://github.com/ggerganov/llama.cpp
!cd llama.cpp && make
!python3 -m pip install -r requirements.txt

!python llama.cpp/convert-hf-to-gguf.py phi2-nl2sql-lora-merged --outfile out/
↪phi2-nl2sql-lora-merged-f16.gguf --outtype f16

```

```

[9]: from zipfile import ZipFile
import os

def get_all_file_paths(directory):
 file_paths = []
 for root, directories, files in os.walk(directory):
 for filename in files:
 filepath = os.path.join(root, filename)
 file_paths.append(filepath)
 return file_paths

def zip_it(directory: str, file_name: str):
 file_paths = get_all_file_paths(directory)
 print('Following files will be zipped:')
 for file_path in file_paths:
 print(file_path)
 with ZipFile(file_name, 'w') as zip:
 for file in file_paths:
 zip.write(file)
 print('All files zipped successfully!')

```

```
[17]: zip_it("phi2-nl2sql-lora-merged", "merged.zip")

from datetime import datetime
import shutil

name = 'merged_' + datetime.now().strftime("%Y_%m_%d_%H_%M_%S") + '.zip'
shutil.move("/content/merged.zip", "/content/drive/MyDrive/phi2_finetune/" +
↳name)
```

Following files will be zipped:

```
pavankumarbalijepalli/phi2-nl2sql-lora-merged/tokenizer_config.json
pavankumarbalijepalli/phi2-nl2sql-lora-merged/tokenizer.json
pavankumarbalijepalli/phi2-nl2sql-lora-merged/model-00001-of-00002.safetensors
pavankumarbalijepalli/phi2-nl2sql-lora-merged/added_tokens.json
pavankumarbalijepalli/phi2-nl2sql-lora-merged/model-00002-of-00002.safetensors
pavankumarbalijepalli/phi2-nl2sql-lora-merged/special_tokens_map.json
pavankumarbalijepalli/phi2-nl2sql-lora-merged/generation_config.json
pavankumarbalijepalli/phi2-nl2sql-lora-merged/model.safetensors.index.json
pavankumarbalijepalli/phi2-nl2sql-lora-merged/config.json
pavankumarbalijepalli/phi2-nl2sql-lora-merged/vocab.json
pavankumarbalijepalli/phi2-nl2sql-lora-merged/merges.txt
All files zipped successfully!
```

```
[17]: '/content/drive/MyDrive/phi2_finetune/merged_2024_02_21_09_19_47.zip'
```

```
[13]: zip_it("out", "phi2-nl2sql-lora-merged-f16.zip")

from datetime import datetime
import shutil

name = 'merged_gguf_' + datetime.now().strftime("%Y_%m_%d_%H_%M_%S") + '.zip'
shutil.move("/content/phi2-nl2sql-lora-merged-f16.zip", "/content/drive/MyDrive/
↳phi2_finetune/" + name)
```

```
[13]: '/content/drive/MyDrive/phi2_finetune/merged_gguf_2024_02_21_10_35_00.zip'
```

```
[]:
```



# eval\_llm

March 1, 2024

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import json
```

```
[2]: # Loading Common Questions and Hard Questions Evaluation Data
phi_eval_c = pd.DataFrame(json.loads(open("./data/phi2_eval_df_c.json", "r").
 ↪read()))
phi_eval_h = pd.DataFrame(json.loads(open("./data/phi2_eval_hard.json", "r").
 ↪read()))
def_eval_c = pd.DataFrame(json.loads(open("./data/sqlc_eval_df_c.json", "r").
 ↪read()))
def_eval_h = pd.DataFrame(json.loads(open("./data/sqlc_eval_hard.json", "r").
 ↪read()))

Loading Training Logs for Fine-tuning Phi-2
training_logs = pd.read_csv("./data/training_logs.csv")

phi_eval_c.shape, phi_eval_h.shape, def_eval_c.shape, def_eval_h.shape, ↵
 ↪training_logs.shape
```

```
[2]: ((1258, 9), (33, 9), (1258, 9), (33, 9), (31, 3))
```

```
[]:
```

```
[3]: phi_eval = pd.concat([phi_eval_c, phi_eval_h])
def_eval = pd.concat([def_eval_c, def_eval_h])

phi_eval.shape, def_eval.shape
```

```
[3]: ((1291, 9), (1291, 9))
```

```
[4]: phi_eval.columns
```

```
[4]: Index(['prompt', 'pred', 'actu', 'inf_time', 'temperature', 'difficulty',
 'token_in', 'token_out', 'tokens_per_sec'],
 dtype='object')
```

```
[5]: phi_eval['tokens_per_sec'] = (phi_eval['token_in'] + phi_eval['token_out'])/
 ↪ phi_eval['inf_time']
 def_eval['tokens_per_sec'] = (def_eval['token_in'] + def_eval['token_out'])/
 ↪ def_eval['inf_time']

 phi_eval['all_tokens'] = phi_eval['token_in'] + phi_eval['token_out']
 def_eval['all_tokens'] = def_eval['token_in'] + def_eval['token_out']

 phi_eval['question'] = phi_eval['prompt'].apply(lambda x :x.split("`")[1])
 def_eval['question'] = def_eval['prompt'].apply(lambda x :x.split("`")[1])

 phi_eval['context'] = phi_eval['prompt'].apply(lambda x :x.split("schema:
 ↪\n")[-1].split("\n\n")[0])
 def_eval['context'] = def_eval['prompt'].apply(lambda x :x.split("schema:
 ↪\n")[-1].split("\n\n")[0])
```

## 0.1 Evaluation Plots

### 0.1.1 Difficulty vs Inference Time

```
[6]: phi_eval.groupby('difficulty')['inf_time'].mean().reset_index().
 ↪ sort_values('inf_time')
```

```
[6]: difficulty inf_time
0 easy 16.261769
2 medium 20.397285
1 hard 38.548714
```

```
[7]: def_eval.groupby('difficulty')['inf_time'].mean().reset_index().
 ↪ sort_values('inf_time')
```


```
[7]: difficulty inf_time
0 easy 30.854192
2 medium 38.260510
1 hard 57.408160
```

```
[8]: # difficulty vs inf_time - bar
figure, axes = plt.subplots(2, 1, sharex=True,
 figsize=(10, 5))
figure.suptitle('Difficulty vs Inference Time - Lower is Better')
figure.tight_layout(pad=3)

axes[0].set_title('Defog Model')
sns.barplot(def_eval, y="difficulty", x="inf_time", ax=axes[0])

axes[1].set_title('Phi-2 Model')
sns.barplot(phi_eval, y="difficulty", x="inf_time", ax=axes[1])
```

```
plt.savefig('./out_images/'+ 'Difficulty vs Inference Time.png',
 bbox_inches='tight')
```

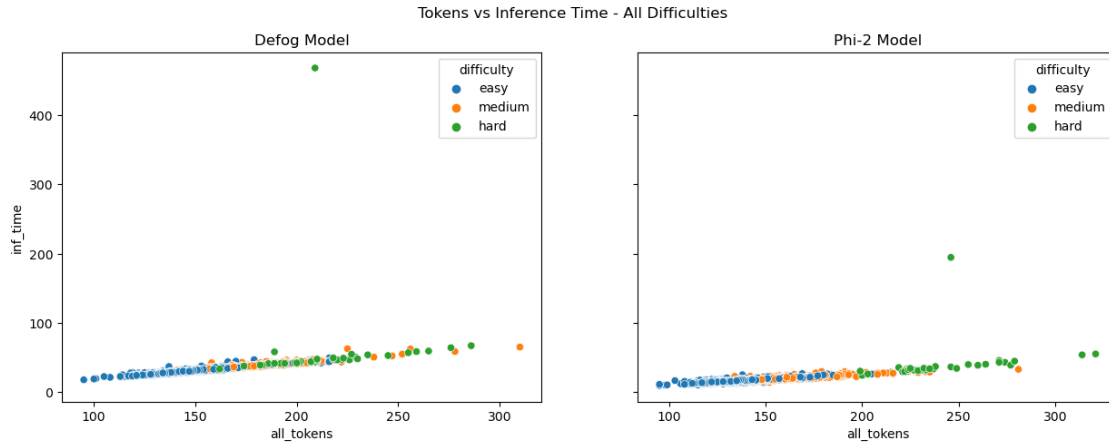


### 0.1.2 Number of Tokens vs Inference Time

```
[9]: # tokens vs inf_time
figure, axes = plt.subplots(1, 2, sharey=True,
 figsize=(15, 5))
figure.suptitle('Tokens vs Inference Time - All Difficulties')

axes[0].set_title('Defog Model')
sns.scatterplot(def_eval, x="all_tokens", y="inf_time", hue="difficulty",
 ax=axes[0])


axes[1].set_title('Phi-2 Model')
sns.scatterplot(phi_eval, x="all_tokens", y="inf_time", hue="difficulty",
 ax=axes[1])
plt.savefig('./out_images/'+ 'Tokens vs Inference Time', bbox_inches='tight')
```



```
[10]: # tokens vs inf_time
figure, axes = plt.subplots(1, 2, sharey=True,
 figsize=(15, 5))
figure.suptitle('Tokens vs Inference Time - Easy and Medium Difficulties')

axes[0].set_title('Defog Model')
sns.scatterplot(def_eval[def_eval['difficulty'].isin(['easy', 'medium'])],
 x="all_tokens", y="inf_time", hue="difficulty", ax=axes[0])

axes[1].set_title('Phi-2 Model')
sns.scatterplot(phi_eval[phi_eval['difficulty'].isin(['easy', 'medium'])],
 x="all_tokens", y="inf_time", hue="difficulty", ax=axes[1])
plt.savefig('./out_images/'+'Tokens vs Inference Time - Easy and Medium Difficulties',
 bbox_inches='tight')
```



### 0.1.3 Number of Tokens per Sec vs Difficulty

```
[11]: phi_eval.groupby('difficulty')['tokens_per_sec'].mean().reset_index().
 ↪sort_values('tokens_per_sec')
```

```
[11]: difficulty tokens_per_sec
1 hard 6.829720
2 medium 8.058363
0 easy 8.390820
```

```
[12]: def_eval.groupby('difficulty')['tokens_per_sec'].mean().reset_index().
 ↪sort_values('tokens_per_sec')
```

```
[12]: difficulty tokens_per_sec
1 hard 4.474975
2 medium 4.726500
0 easy 4.854409
```


```
[45]: round((8.390820-4.854409)/8.390820,5)*100, round((8.058363-4.726500)/8.
 ↪058363,5)*100, round((6.829720-4.474975)/6.829720,5)*100
```

```
[45]: (42.146, 41.347, 34.477999999999994)
```

```
[13]: # difficulty vs tokens_per_inf_time - bar
figure, axes = plt.subplots(2, 1, sharex=True,
 figsize=(10, 5))
figure.suptitle('Difficulty vs Tokens per Sec - Higher is Better')
figure.tight_layout(pad=3)

axes[0].set_title('Defog Model')
sns.barplot(def_eval.groupby('difficulty')['tokens_per_sec'].mean().
 ↪reset_index().sort_values('tokens_per_sec'),y="difficulty",
 ↪x="tokens_per_sec", ax=axes[0])

axes[1].set_title('Phi-2 Model')
sns.barplot(phi_eval.groupby('difficulty')['tokens_per_sec'].mean().
 ↪reset_index().sort_values('tokens_per_sec'),y="difficulty",
 ↪x="tokens_per_sec", ax=axes[1])
plt.savefig('./out_images/'+ 'Difficulty vs Tokens per Sec', bbox_inches='tight')
```



#### 0.1.4 Tokens Generated vs Difficulty

```
[47]: phi_eval.groupby('difficulty')['token_out'].mean().reset_index().
 ↪sort_values('token_out')
```

```
[47]: difficulty token_out
0 easy 25.839450
2 medium 33.673684
1 hard 81.692308
```

```
[48]: def_eval.groupby('difficulty')['token_out'].mean().reset_index().
 ↪sort_values('token_out')
```

```
[48]: difficulty token_out
0 easy 34.969037
2 medium 46.142105
1 hard 65.538462
```

```
[49]: round((34.969037-25.839450)/34.969037,5)*100, round((46.142105-33.673684)/46.
 ↪142105,5)*100, round((65.538462-81.692308)/81.692308,5)*100
```

```
[49]: (26.107999999999997, 27.022000000000002, -19.774)
```


```
[14]: # difficulty vs out tokens - bar
figure, axes = plt.subplots(2, 1, sharex=True,
 figsize=(10, 5))
figure.suptitle('Difficulty vs Out Tokens - Higher is Better')
figure.tight_layout(pad=3)
```

```

axes[0].set_title('Defog Model')
sns.barplot(def_eval, y="difficulty", x="token_out", ax=axes[0])

axes[1].set_title('Phi-2 Model')
sns.barplot(phi_eval, y="difficulty", x="token_out", ax=axes[1])
plt.savefig('./out_images/'+ 'Difficulty vs Out Tokens', bbox_inches='tight')

```



### 0.1.5 Execution Success

```

[69]: import sqlite3

Create DB schema based on context and run the query to check for execution.
def execution_test(row):
 conn = sqlite3.connect(':memory:')
 script = row['context']
 success = True
 try:
 conn.executescript(script)
 sql_query = pd.read_sql_query(row['pred'], conn)
 except:
 success = False
 return success

phi_eval['exec_success'] = phi_eval.apply(lambda x: execution_test(x), axis=1).
 ↪map({True: '1', False: '0'})
def_eval['exec_success'] = def_eval.apply(lambda x: execution_test(x), axis=1).
 ↪map({True: '1', False: '0'})

```

```
phi_eval['exec_success_int'] = phi_eval.apply(lambda x: execution_test(x),
↪axis=1).map({True: 1, False: 0})
def_eval['exec_success_int'] = def_eval.apply(lambda x: execution_test(x),
↪axis=1).map({True: 1, False: 0})
```

```
[75]: phi_eval.groupby('difficulty')['exec_success_int'].mean().reset_index().
↪sort_values('exec_success_int'), def_eval.
↪groupby('difficulty')['exec_success_int'].mean().reset_index().
↪sort_values('exec_success_int')
```

```
[75]: (difficulty exec_success_int
1 hard 0.487179
2 medium 0.918421
0 easy 0.985092,
 difficulty exec_success_int
2 medium 0.513158
1 hard 0.564103
0 easy 0.620413)
```

```
[77]: round((22-19)/22,5)*100, round((349-195)/349,5)*100, round((859-541)/
↪859,5)*100, (98+91+48)/300, (62+56+51)/300
```

```
[77]: (13.636000000000001, 44.126, 37.019999999999996, 0.79, 0.5633333333333334)
```

```
[25]: # difficulty vs inf_time - bar
figure, axes = plt.subplots(1, 2, sharex=True,
 figsize=(10, 5))
figure.suptitle('Difficulty vs Execution Success - Higher 1s are Better')
figure.tight_layout(pad=3)

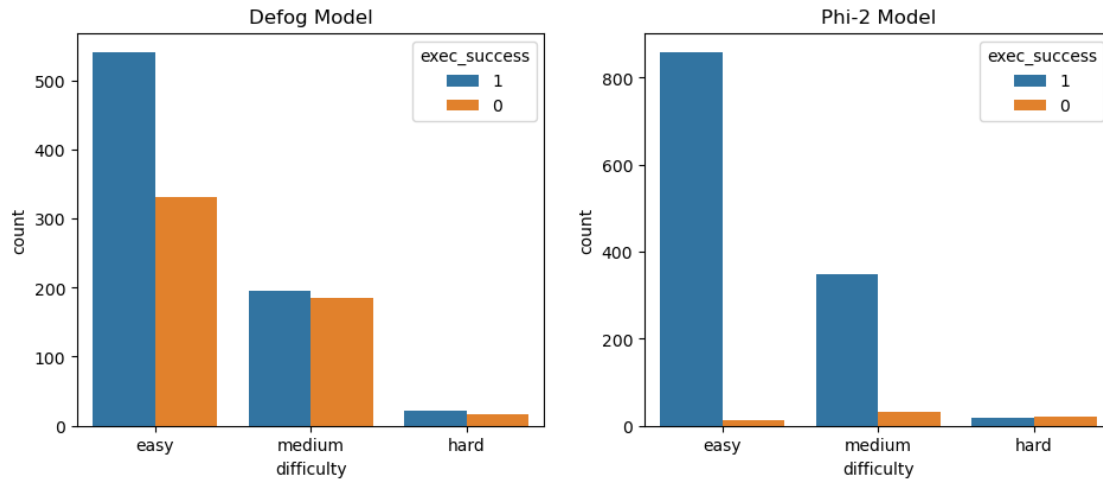
axes[0].set_title('Defog Model')
sns.countplot(def_eval, x="difficulty", ax=axes[0], hue="exec_success",
↪hue_order="10")

axes[1].set_title('Phi-2 Model')
sns.countplot(phi_eval, x="difficulty", ax=axes[1], hue="exec_success")

plt.savefig('./out_images/'+'Difficulty vs Execution Success',
↪bbox_inches='tight')
```



Difficulty vs Execution Success - Higher 1s are Better



### 0.1.6 Exact Matches

```
[17]: ## Check perfect matches based on the actual value
phi_eval['exact_match'] = phi_eval['pred'].str.lower().str.strip() == \
 ↪ phi_eval['actu'].str.lower().str.strip()
def_eval['exact_match'] = def_eval['pred'].str.lower().str.strip() == \
 ↪ def_eval['actu'].str.lower().str.strip()

phi_eval['exact_match'] = phi_eval['exact_match'].map({True: '1', False: '0'})
def_eval['exact_match'] = def_eval['exact_match'].map({True: '1', False: '0'})
```

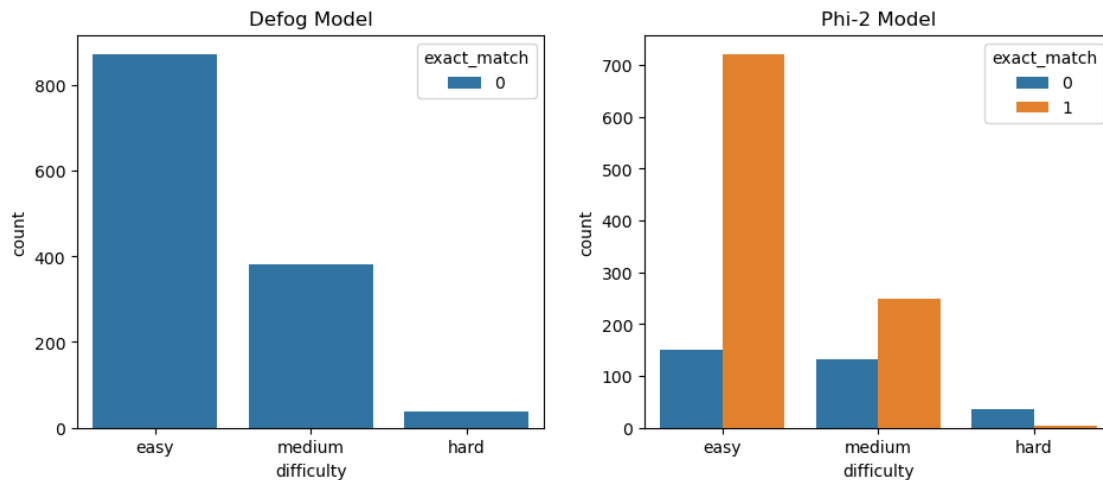
```
[18]: # difficulty vs inf_time - bar
figure, axes = plt.subplots(1, 2, sharex=True, figsize=(10, 5))
figure.suptitle('Difficulty vs Exact Match - Higher 1s is Better')
figure.tight_layout(pad=3)

axes[0].set_title('Defog Model')
sns.countplot(def_eval, x="difficulty", ax=axes[0], hue="exact_match")

axes[1].set_title('Phi-2 Model')
sns.countplot(phi_eval, x="difficulty", ax=axes[1], hue="exact_match")

plt.savefig('./out_images/' + 'Difficulty vs Exact Match', bbox_inches='tight')
```

Difficulty vs Exact Match - Higher 1s is Better



### 0.1.7 Training Logs

```
[19]: training_logs['train_accuracy'] = 1 - training_logs['Training Loss']
 training_logs['val_accuracy'] = 1 - training_logs['Validation Loss']

[20]: training_logs.columns = ['step', 'train_loss', 'val_loss', 'train_accuracy', '
 ↪ val_accuracy']

[21]: # training_logs
 figure, axes = plt.subplots(1, 1, sharex=True, figsize=(10, 5))
 figure.suptitle('Training and Validation Accuracy vs Steps')
 sns.lineplot(data=training_logs, y="train_accuracy", x="step", label="train_
 ↪ accuracy")
 sns.lineplot(data=training_logs, y="val_accuracy", x="step", label="validation_
 ↪ accuracy")

 plt.savefig('./out_images/'+ 'Training and Validation Accuracy vs Steps',
 ↪ bbox_inches='tight')
```

```
C:\Users\pavan\anaconda3\envs\phi2\Lib\site-packages\seaborn_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
 with pd.option_context('mode.use_inf_as_na', True):
C:\Users\pavan\anaconda3\envs\phi2\Lib\site-packages\seaborn_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
 with pd.option_context('mode.use_inf_as_na', True):
C:\Users\pavan\anaconda3\envs\phi2\Lib\site-packages\seaborn_oldcore.py:1119:
```

FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.


```
with pd.option_context('mode.use_inf_as_na', True):
```

C:\Users\pavan\anaconda3\envs\phi2\Lib\site-packages\seaborn\\_oldcore.py:1119:

FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```

Training and Validation Accuracy vs Steps



```
[22]: # training_logs
figure, axes = plt.subplots(1, 1, sharex=True, figsize=(10, 5))
figure.suptitle('Training and Validation Loss vs Steps')
sns.lineplot(data=training_logs, y="train_loss", x="step", label="train loss")
sns.lineplot(data=training_logs, y="val_loss", x="step", label="validation_
↳ loss")

plt.savefig('./out_images/'+'Training and Validation Loss vs Steps',
↳ bbox_inches='tight')
```

C:\Users\pavan\anaconda3\envs\phi2\Lib\site-packages\seaborn\\_oldcore.py:1119:

FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```

C:\Users\pavan\anaconda3\envs\phi2\Lib\site-packages\seaborn\\_oldcore.py:1119:

FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```

C:\Users\pavan\anaconda3\envs\phi2\Lib\site-packages\seaborn\\_oldcore.py:1119:

FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.


```
with pd.option_context('mode.use_inf_as_na', True):
```

C:\Users\pavan\anaconda3\envs\phi2\Lib\site-packages\seaborn\\_oldcore.py:1119:

FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.


```
with pd.option_context('mode.use_inf_as_na', True):
```

Training and Validation Loss vs Steps




```
[23]: shapes = pd.DataFrame({"split": ["training", "testing", "evaluation"], "size": [63647, 7858, 7072]}, index=[1, 2, 3])
sns.barplot(shapes, y="split", x="size")
```

```
[23]: <Axes: xlabel='size', ylabel='split'>
```



```
[35]: usage = pd.DataFrame({"split": ["phi-2", "defog"], "size": [11.2, 15.5]},
 ↪ index=[1, 2])
sns.barplot(usage, x="split", y="size",width=0.4)
```

```
[35]: <Axes: xlabel='split', ylabel='size'>
```



```
[]:
```